



*cassandra*

# 从零开始Cassandra

应用，原理，选型

李力

2022年11月16日

Code: <https://github.com/xiaozhiliaoo/cassandra-practice>

Slide: <https://github.com/xiaozhiliaoo/my-slides/blob/master/Database-Cassandra.pptx>

1

## 简介

Cassandra是什么？

2

## 数据建模与开发

项目使用指南

3

## 系统架构

Cassandra单机，分布式原理

4

## NoSQL存储选型

我们应该关注什么？

5

## 参考资料

基础与进阶

6

## 回顾&QA

思维导图

# 1. 简介

The content of the title

- Apache Cassandra is an open source, distributed, **shared-nothing, decentralized (P2P对等结构化覆盖网络)**, elastically scalable, highly available, fault-tolerant, **tuneably consistent (C“A”P, PACELC)**, **row-oriented (宽列)** database that bases its distribution design on Amazon's **Dynamo** and its data model on Google's **Bigtable**. Created at Facebook, it is now used at some of the most popular sites on the Web.
- Apache Cassandra is an open source NoSQL **distributed database** trusted by thousands of companies for scalability and high availability without compromising performance. **Linear scalability** and proven fault-tolerance on **commodity hardware** or cloud infrastructure make it the perfect platform for **mission-critical** data.
- Cassandra World Party

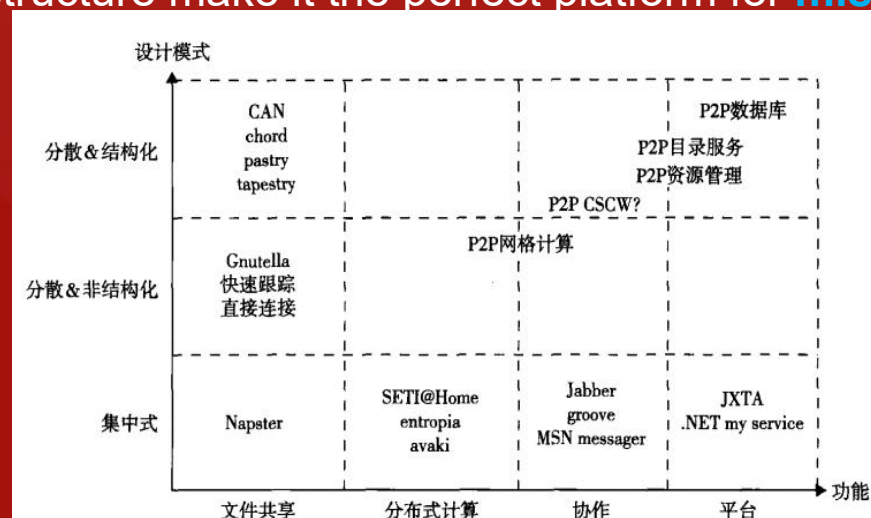


图 8-9 按照功能和设计模式对 P2P 系统进行分类

# 1. 简介-db-engines排名

The content of the title

| 397 systems in ranking, November 2022 |          |          |                                   |                            |          |          |          |  |  |
|---------------------------------------|----------|----------|-----------------------------------|----------------------------|----------|----------|----------|--|--|
| Rank                                  |          |          | DBMS                              | Database Model             | Score    |          |          |  |  |
| Nov 2022                              | Oct 2022 | Nov 2021 |                                   |                            | Nov 2022 | Oct 2022 | Nov 2021 |  |  |
| 1.                                    | 1.       | 1.       | Oracle                            | Relational, Multi-model    | 1241.69  | +5.32    | -31.04   |  |  |
| 2.                                    | 2.       | 2.       | MySQL                             | Relational, Multi-model    | 1205.54  | +0.17    | -5.98    |  |  |
| 3.                                    | 3.       | 3.       | Microsoft SQL Server              | Relational, Multi-model    | 912.51   | -12.17   | -41.78   |  |  |
| 4.                                    | 4.       | 4.       | PostgreSQL                        | Relational, Multi-model    | 623.16   | +0.44    | +25.88   |  |  |
| 5.                                    | 5.       | 5.       | MongoDB                           | Document, Multi-model      | 477.90   | -8.33    | -9.45    |  |  |
| 6.                                    | 6.       | 6.       | Redis                             | Key-value, Multi-model     | 182.05   | -1.33    | +10.55   |  |  |
| 7.                                    | 7.       | 8.       | Elasticsearch                     | Search engine, Multi-model | 150.32   | -0.74    | -8.76    |  |  |
| 8.                                    | 8.       | 7.       | IBM Db2                           | Relational, Multi-model    | 149.56   | -0.10    | -17.96   |  |  |
| 9.                                    | 9.       | 11.      | Microsoft Access                  | Relational                 | 135.03   | -3.14    | +15.79   |  |  |
| 10.                                   | 10.      | 9.       | SQLite                            | Relational                 | 134.63   | -3.17    | +4.83    |  |  |
| 11.                                   | 11.      | 10.      | Cassandra                         | Wide column                | 118.12   | +0.18    | -2.76    |  |  |
| 12.                                   | 13.      | 18.      | Snowflake                         | Relational                 | 110.15   | +3.43    | +45.97   |  |  |
| 13.                                   | 12.      | 12.      | MariaDB                           | Relational, Multi-model    | 104.91   | -4.40    | +2.72    |  |  |
| 14.                                   | 14.      | 13.      | Splunk                            | Search engine              | 94.23    | -0.43    | +1.92    |  |  |
| 15.                                   | 15.      | 16.      | Amazon DynamoDB                   | Multi-model                | 85.40    | -2.95    | +8.41    |  |  |
| 16.                                   | 16.      | 15.      | Microsoft Azure SQL Database      | Relational, Multi-model    | 83.66    | -1.30    | +2.34    |  |  |
| 17.                                   | 17.      | 14.      | Hive                              | Relational                 | 81.89    | +1.29    | -1.42    |  |  |
| 18.                                   | 18.      | 17.      | Teradata                          | Relational, Multi-model    | 65.23    | -0.84    | -4.35    |  |  |
| 19.                                   | 20.      |          | Databricks                        | Multi-model                | 60.89    | +3.28    |          |  |  |
| 20.                                   | 19.      | 19.      | Neo4j                             | Graph                      | 57.30    | -1.38    | -0.68    |  |  |
| 21.                                   | 23.      | 21.      | FileMaker                         | Relational                 | 54.31    | +1.90    | +0.08    |  |  |
| 22.                                   | 22.      | 24.      | Google BigQuery                   | Relational                 | 54.13    | +1.68    | +9.13    |  |  |
| 23.                                   | 24.      | 20.      | SAP HANA                          | Relational, Multi-model    | 51.45    | -0.63    | -4.08    |  |  |
| 24.                                   | 21.      | 22.      | Solr                              | Search engine, Multi-model | 51.33    | -2.17    | -2.52    |  |  |
| 25.                                   | 25.      | 23.      | SAP Adaptive Server               | Relational, Multi-model    | 43.58    | +0.62    | -7.35    |  |  |
| 26.                                   | 26.      | 25.      | HBase                             | Wide column                | 40.41    | -1.25    | -4.59    |  |  |
| 27.                                   | 27.      | 26.      | Microsoft Azure Cosmos DB         | Multi-model                | 39.75    | -0.67    | -1.08    |  |  |
| 28.                                   | 28.      | 27.      | PostGIS                           | Spatial DBMS, Multi-model  | 30.78    | -0.08    | -1.14    |  |  |
| 29.                                   | 29.      | 29.      | InfluxDB                          | Time Series, Multi-model   | 29.96    | +0.38    | +1.42    |  |  |
| 30.                                   | 30.      | 28.      | Couchbase                         | Document, Multi-model      | 28.62    | +0.66    | -1.25    |  |  |
| 31.                                   | 31.      | 32.      | Amazon Redshift                   | Relational                 | 27.04    | +0.03    | +2.17    |  |  |
| 32.                                   | 33.      | 30.      | Firebird                          | Relational                 | 25.37    | +0.36    | -1.54    |  |  |
| 33.                                   | 32.      | 31.      | Memcached                         | Key-value                  | 24.45    | -0.95    | -1.92    |  |  |
| 34.                                   | 34.      | 39.      | Microsoft Azure Synapse Analytics | Relational                 | 23.03    | -0.07    | +5.13    |  |  |
| 35.                                   | 35.      | 34.      | Informix                          | Relational, Multi-model    | 22.82    | -0.05    | -0.41    |  |  |
| 36.                                   | 36.      | 33.      | Spark SQL                         | Relational                 | 21.90    | -0.68    | -1.43    |  |  |
| 37.                                   | 37.      | 36.      | Firebase Realtime Database        | Document                   | 19.84    | -0.29    | -0.23    |  |  |
| 38.                                   | 38.      | 35.      | Vertica                           | Relational, Multi-model    | 19.70    | -0.15    | -0.63    |  |  |
| 39.                                   | 40.      | 38.      | Impala                            | Relational, Multi-model    | 17.92    | +0.24    | -1.01    |  |  |
| 40.                                   | 39.      | 37.      | Netezza                           | Relational                 | 17.27    | -0.57    | -2.52    |  |  |
| 41.                                   | 41.      | 40.      | CouchDB                           | Document, Multi-model      | 15.98    | -0.27    | -0.82    |  |  |

| Rank     |          |          | DBMS                          | Database Model           | Score    |          |          |
|----------|----------|----------|-------------------------------|--------------------------|----------|----------|----------|
| Nov 2022 | Oct 2022 | Nov 2021 |                               |                          | Nov 2022 | Oct 2022 | Nov 2021 |
| 1.       | 1.       | 1.       | Cassandra                     | Wide column              | 118.12   | +0.18    | -2.76    |
| 2.       | 2.       | 2.       | HBase                         | Wide column              | 40.41    | -1.25    | -4.59    |
| 3.       | 3.       | 3.       | Microsoft Azure Cosmos DB     | Multi-model              | 39.75    | -0.67    | -1.08    |
| 4.       | 4.       | 4.       | Datastax Enterprise           | Wide column, Multi-model | 8.68     | +0.39    | -1.01    |
| 5.       | 5.       | 5.       | Microsoft Azure Table Storage | Wide column              | 6.02     | +0.05    | +0.33    |
| 6.       | 6.       | 6.       | Google Cloud Bigtable         | Multi-model              | 5.24     | -0.20    | +0.97    |
| 7.       | 8.       | 8.       | Accumulo                      | Wide column              | 4.98     | +0.41    | +1.07    |
| 8.       | 7.       | 7.       | ScyllaDB                      | Wide column, Multi-model | 4.95     | +0.26    | +1.01    |
| 9.       | 9.       | 9.       | HPE Ezmeral Data Fabric       | Multi-model              | 1.26     | +0.13    | +0.37    |
| 10.      | 10.      | 11.      | Amazon Keyspaces              | Wide column              | 0.77     | -0.02    | +0.24    |
| 11.      | 11.      | 10.      | Elassandra                    | Wide column, Multi-model | 0.57     | +0.09    | -0.06    |
| 12.      | 12.      | 12.      | Alibaba Cloud Table Store     | Wide column              | 0.33     | -0.03    | -0.08    |
| 13.      | 13.      | 13.      | SWC-DB                        | Wide column, Multi-model | 0.07     | -0.02    | +0.07    |



# 1. 简介

The content of the title

## Bigtable: A Distributed Storage System for Structured Data

2006

Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach  
Mike Burrows, Tushar Chandra, Andrew Fikes, Robert E. Gruber

{fay,jeff,sanjay,wilsonh,kerr,m3b,tushar,fikes,gruber}@google.com

Google, Inc.

## 2007 Dynamo: Amazon's Highly Available Key-value Store

Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati,  
Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall  
and Werner Vogels

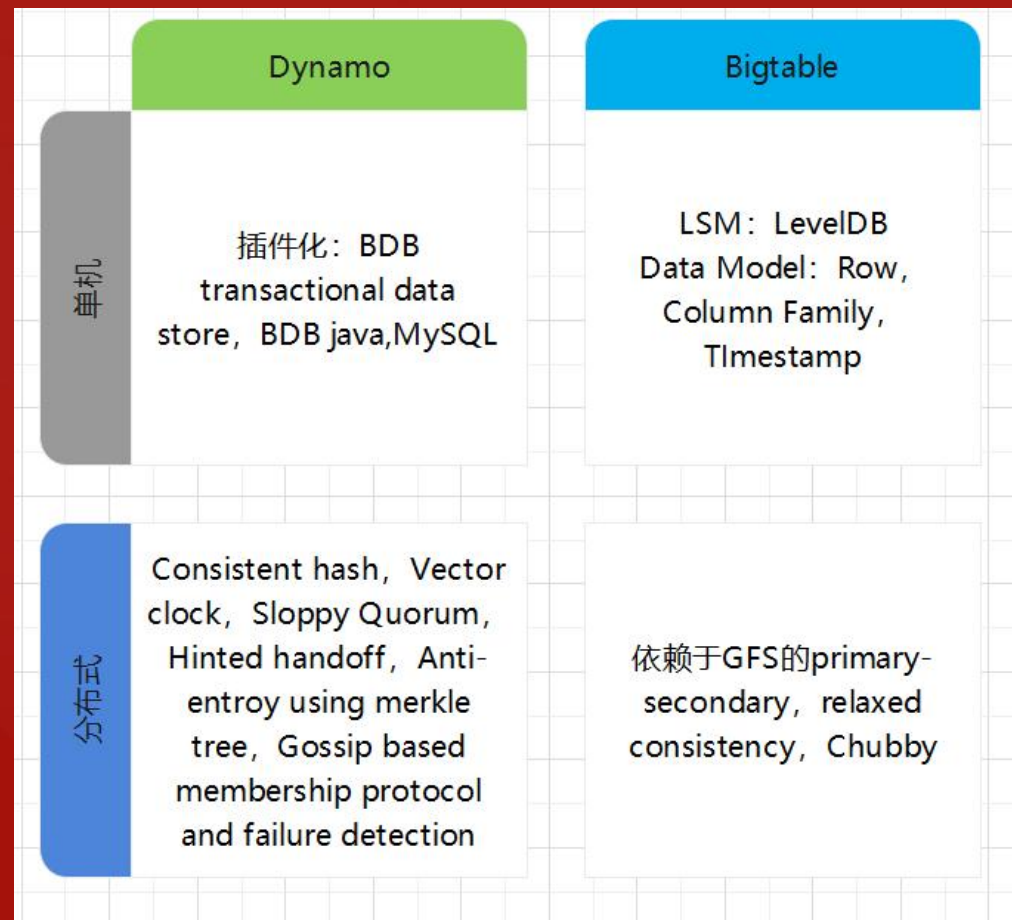
Amazon.com

2009

## Cassandra - A Decentralized Structured Storage System

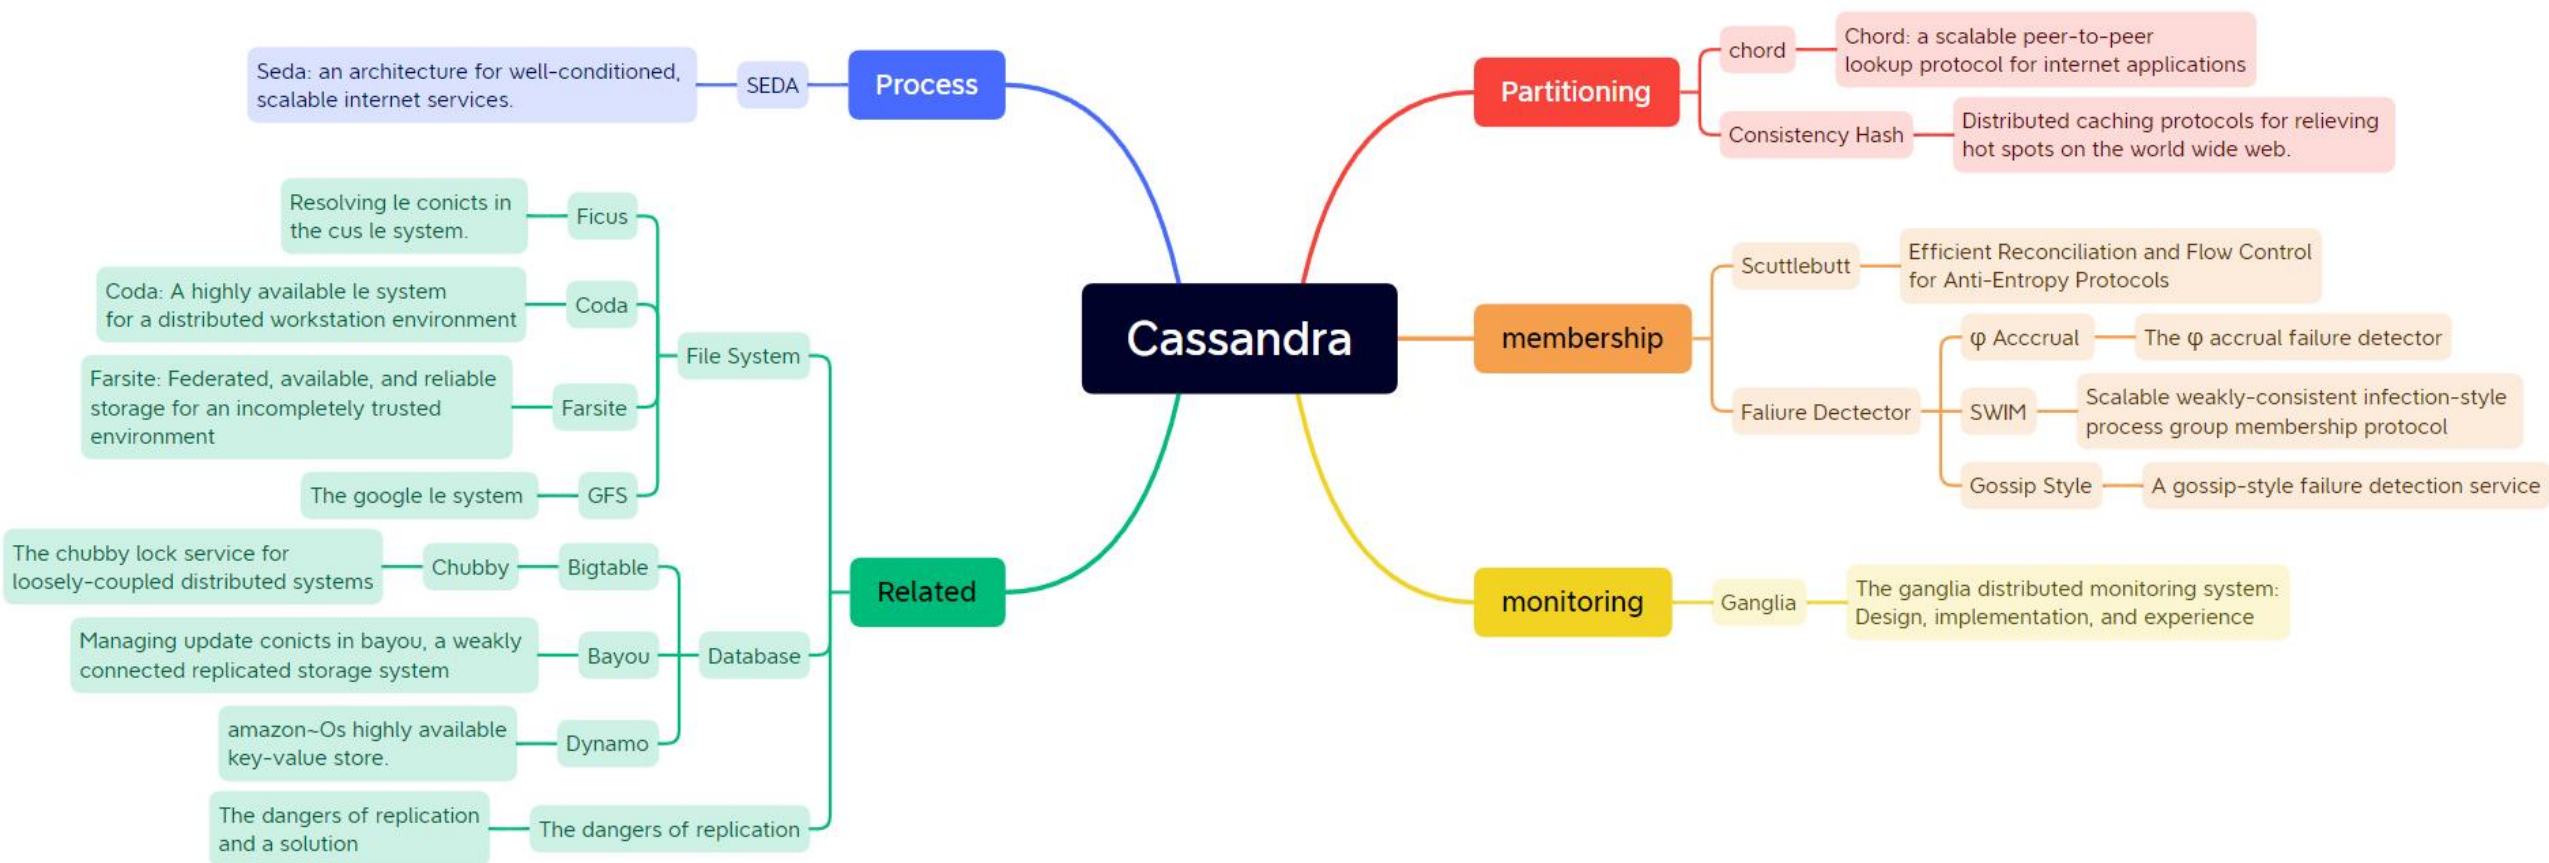
Avinash Lakshman  
Facebook

Prashant Malik  
Facebook



# 1. 简介

The content of the title

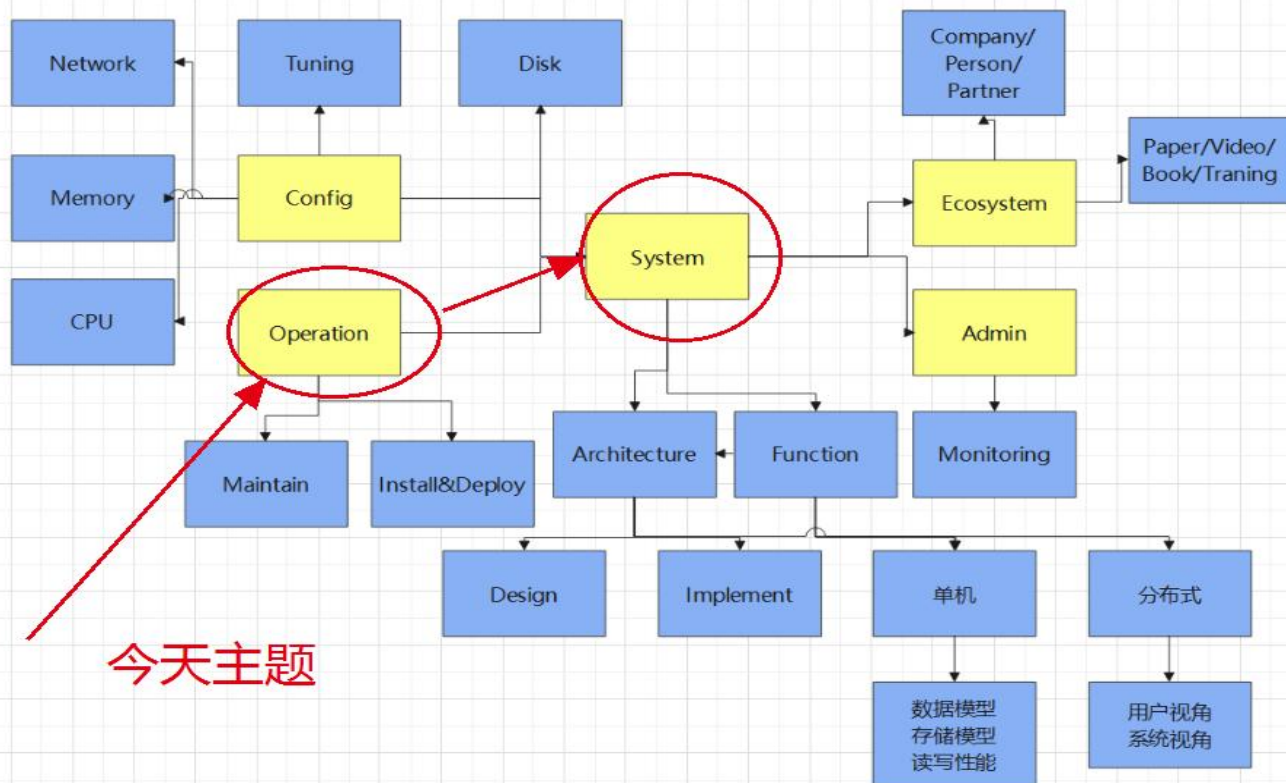


# 1. 简介-研究模型

The content of the title

## • COSEA模型

### • 参考链接



## Main documentation

|                                |  |
|--------------------------------|--|
| Getting started                | Newbie starting point                    |
| What's new in 4.1              | What's new in Cassandra 4.1              |
| Architecture                   | Cassandra's big picture                  |
| Data modeling                  | Hint: it's not relational                |
| Cassandra Query Language (CQL) | CQL reference documentation              |
| Configuration                  | Cassandra's handles and knobs            |
| Operation                      | The operator's corner                    |
| Tools                          | cqlsh, nodetool, and others              |
| Troubleshooting                | What to look for when you have a problem |
| FAQ                            | Frequently asked questions               |
| Plug-ins                       | Third-party plug-ins                     |
| Native Protocols               | Native Cassandra protocol specifications |

## 2. 数据建模-概念

The content of the title

---

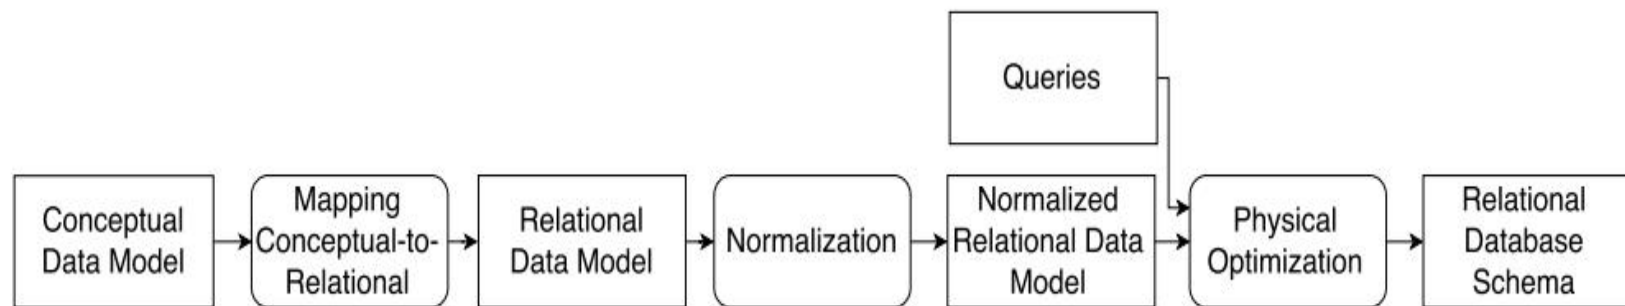
- 理论
  - 范式和反范式，数仓（Star/Snowflake），可变和不可变
- 产品
  - 网状，关系型，NoSQL(KV,文档,宽列,图形)，搜索，数仓列式
- 案例
  - 预定酒店建模
- 重要
  - 模型影响数据存储以及查询。




## 2. 数据建模-过程

The content of the title

- 建模过程：查询驱动建模



(a) Relational Data Modeling

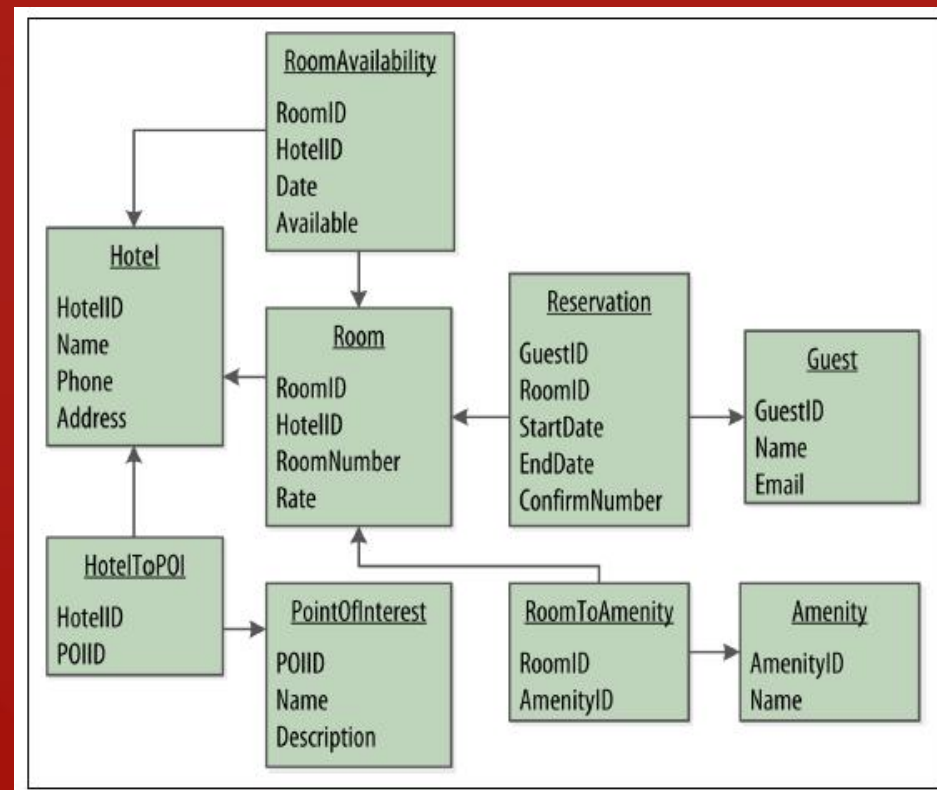
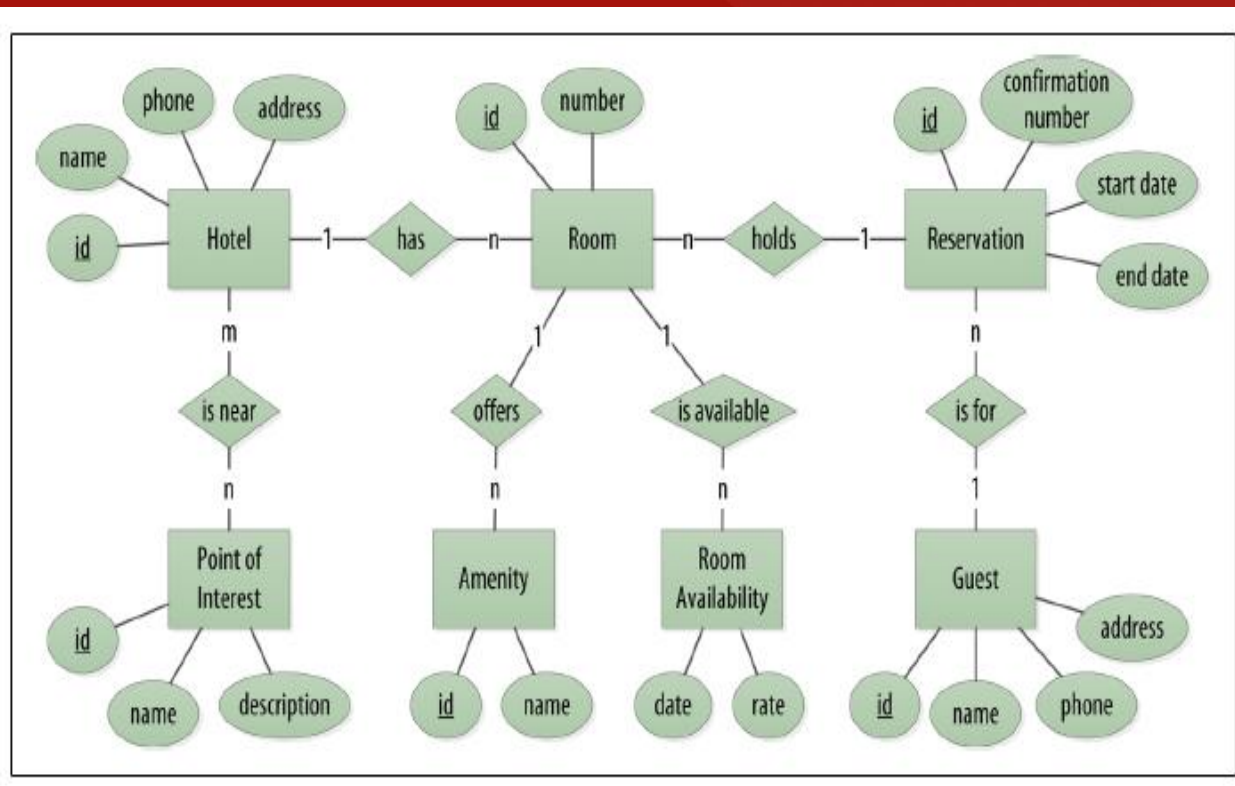


(b) Cassandra Data Modeling

## 2. 数据建模-案例

The content of the title

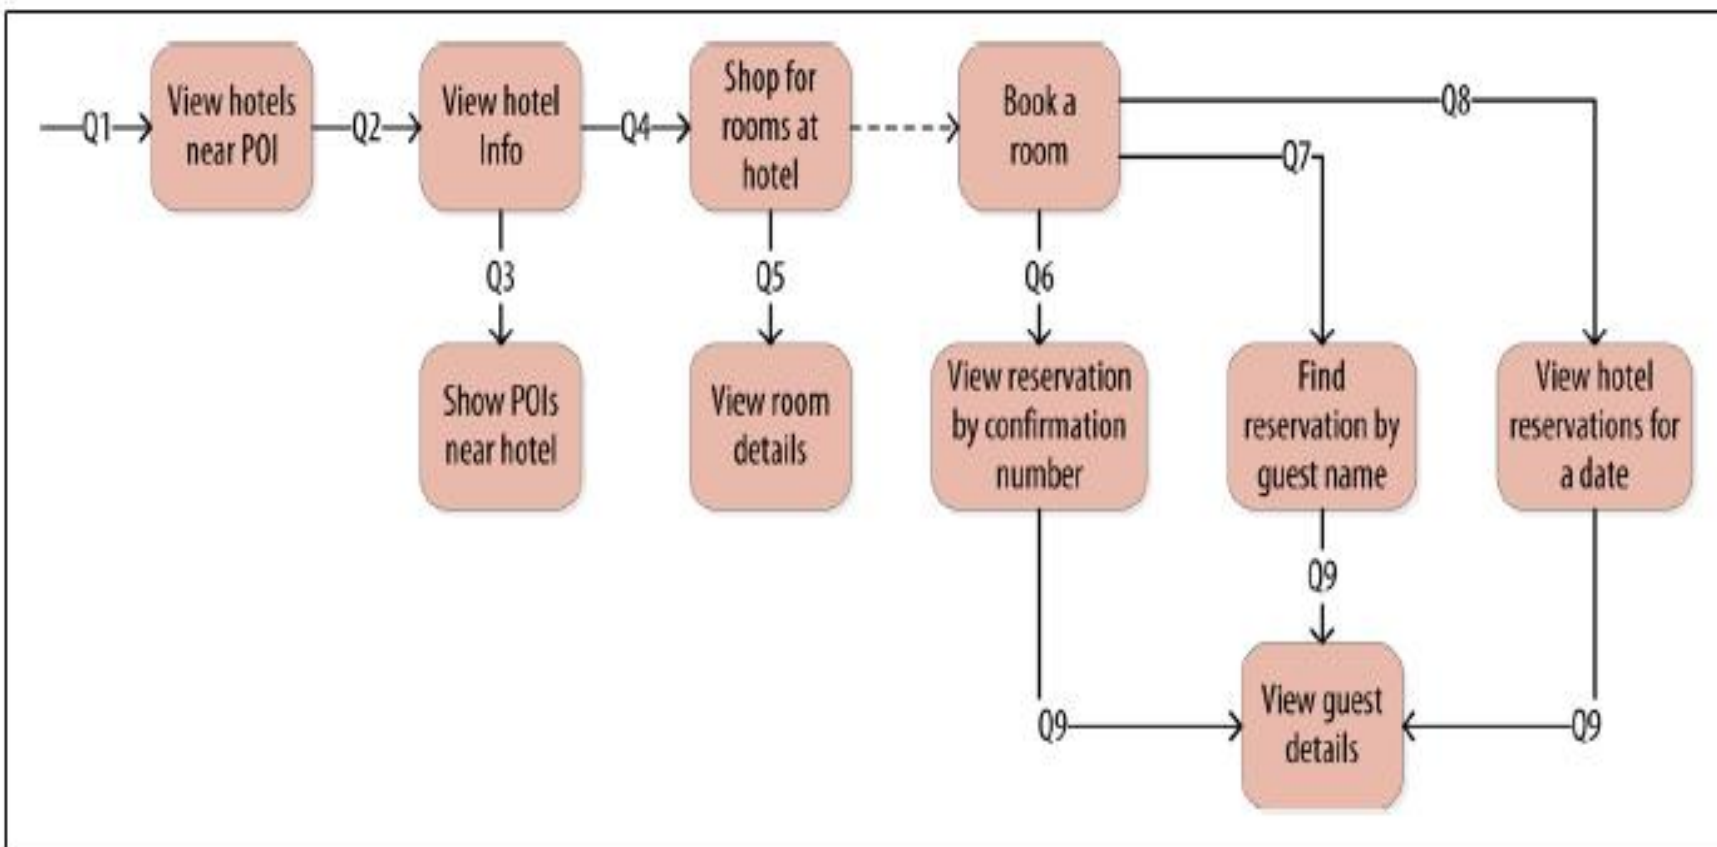
### • 酒店预定



## 2. 数据建模-案例- workflow

The content of the title

### • 定义应用查询



## 2. 数据建模-案例-逻辑数据建模

The content of the title

- 逻辑数据建模
  - 酒店领域
  - 预定领域

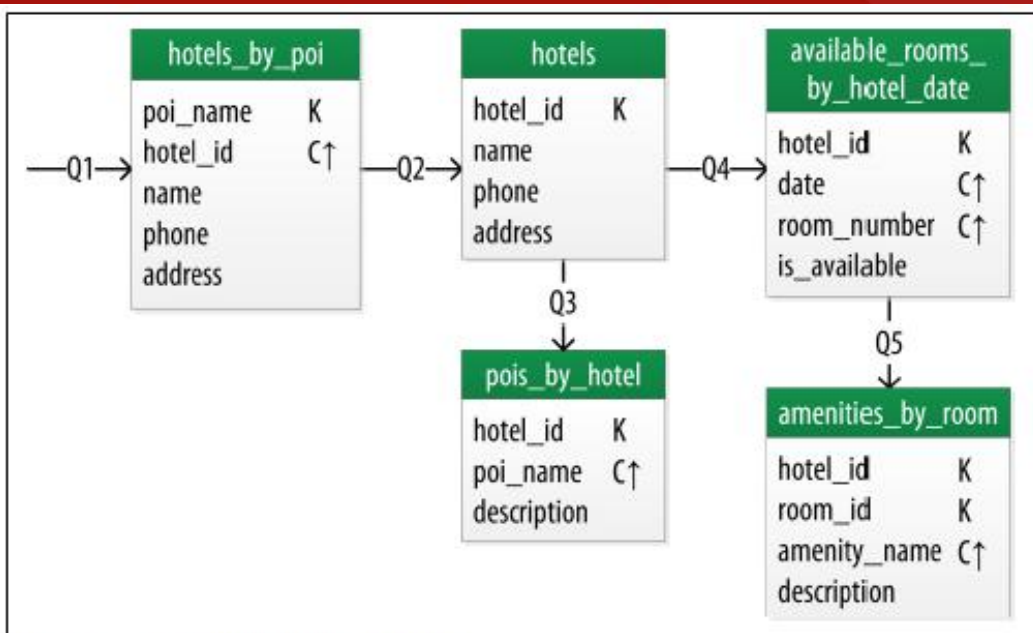


Figure 5-5. Hotel domain logical model

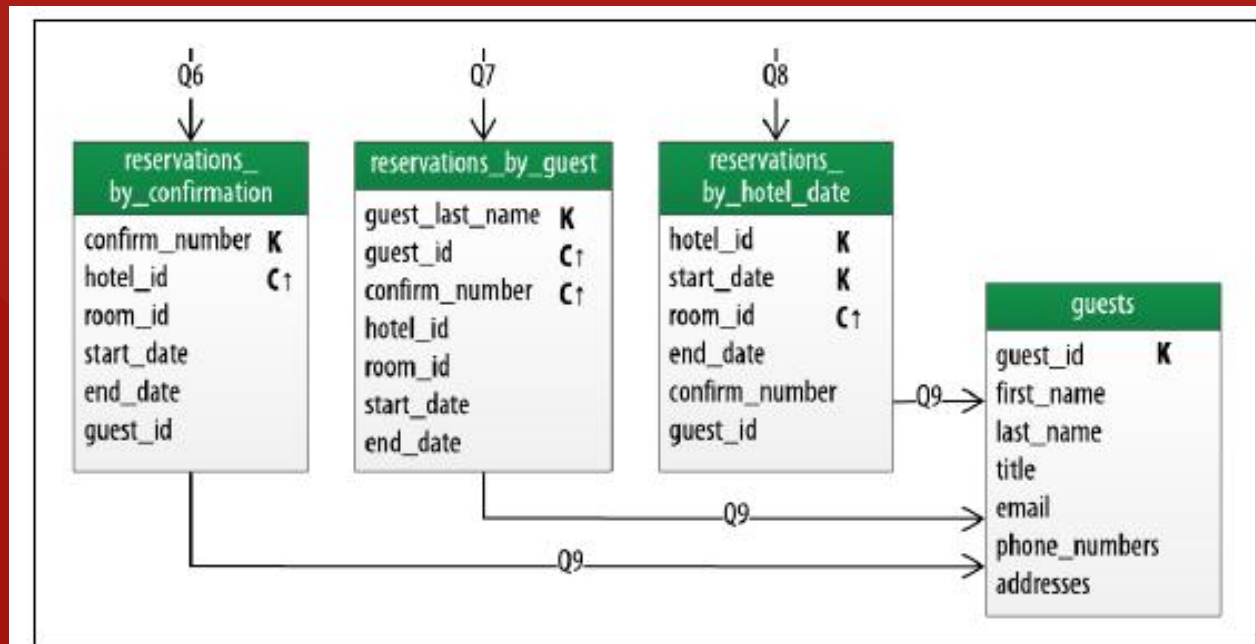


Figure 5-6. A denormalized logical model for reservations



## 2. 数据建模-案例-物理数据建模

The content of the title

- 物理数据建模（Chebotko图）
  - 酒店keyspace
  - 预定keyspace

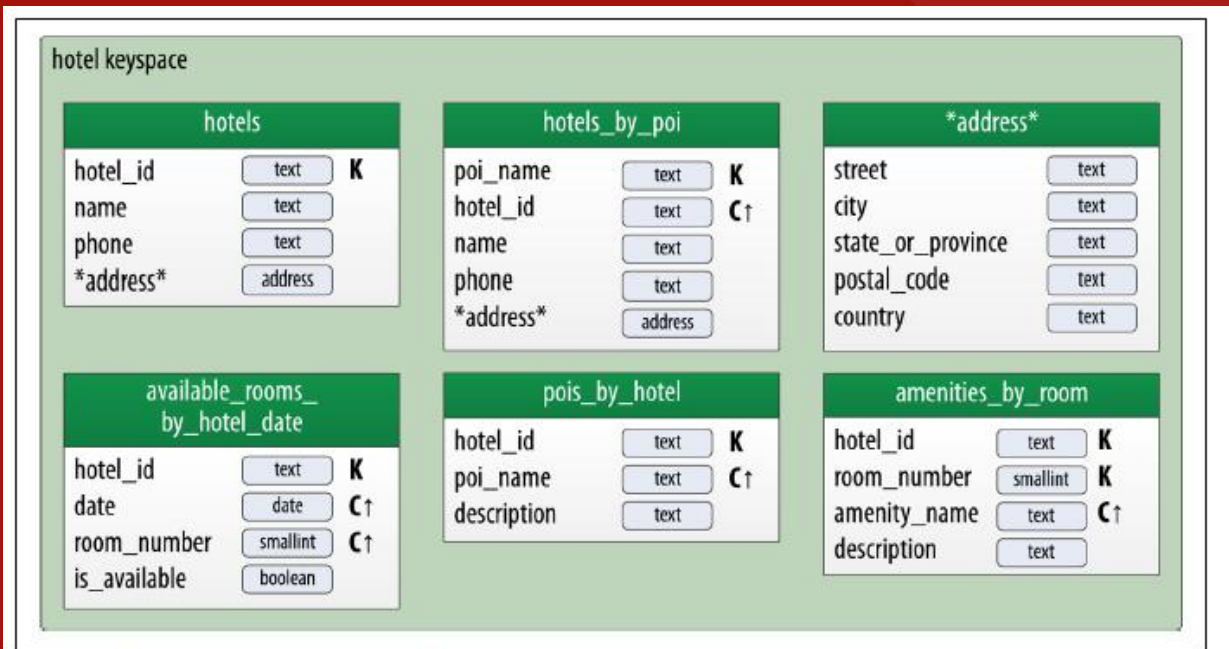


Figure 5-8. Hotel physical model

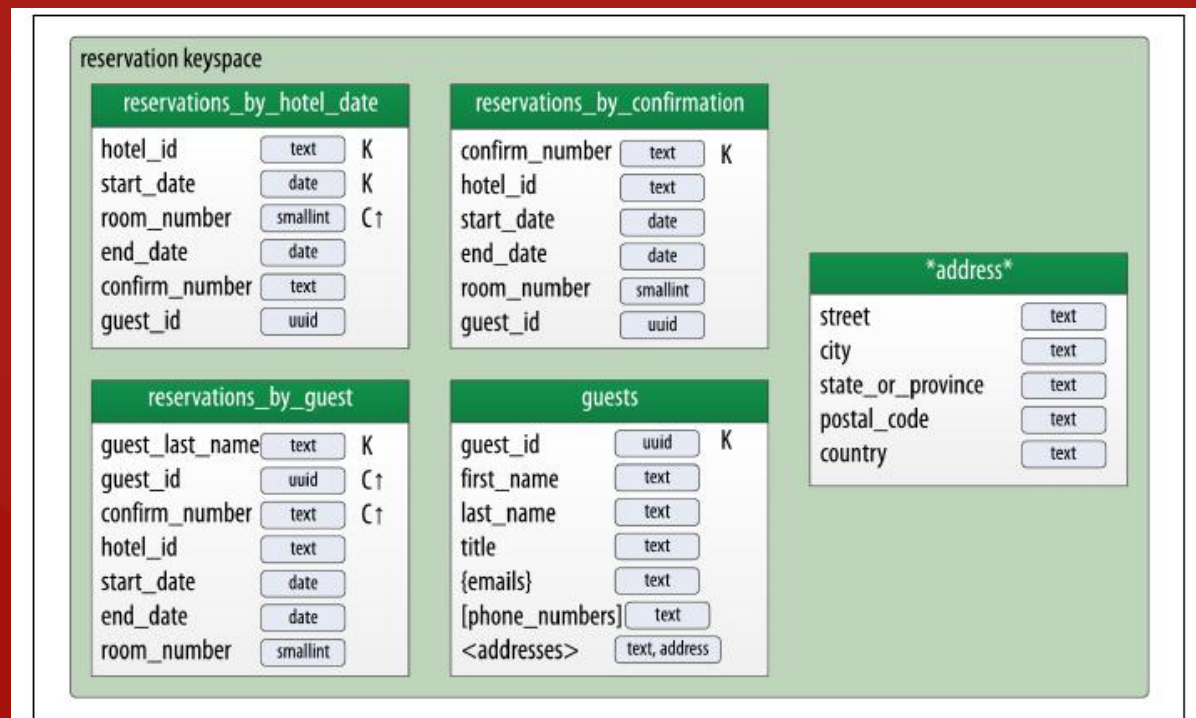


Figure 5-9. Reservation physical model

## 2. 数据建模-案例-评估和改进

The content of the title

- 计算分区大小

$$N_v = N_r(N_c - N_{pk} - N_s) + N_s$$

- 计算磁盘大小

$$S_t = \sum_i \text{sizeOf}(c_{k_i}) + \sum_j \text{sizeOf}(c_{s_j}) + N_r \times \left( \sum_k \text{sizeOf}(c_{r_k}) + \sum_l \text{sizeOf}(c_{c_l}) \right) + N_v \times \text{sizeOf}(t_{avg})$$

- 分解大分区

- 分区键加列，分桶

| available_rooms_<br>by_hotel_date |          |    |
|-----------------------------------|----------|----|
| hotel_id                          | text     | K  |
| date                              | date     | C↑ |
| room_number                       | smallint | C↑ |
| is_available                      | boolean  |    |

## 2. 数据建模-案例-定义数据Schema

The content of the title

---

- 酒店：[参考](#)
- 预定：[参考](#)

## 2. 数据建模-案例-总结

The content of the title

---

- 反范式化
- 查询驱动建模
- 主键查询优先
- 搜索/聚合转es，分析spark或数仓
- 常见错误整理：[参考](#)



## 2.2. 数据建模-客户端

The content of the title

---

- 线上服务器版本 [cqlsh 5.0.1 | Cassandra 3.11.4 | CQL spec 3.4.4 | Native protocol v4]
- Driver版本: 3.7.0

## 2. 数据建模-客户端

The content of the title

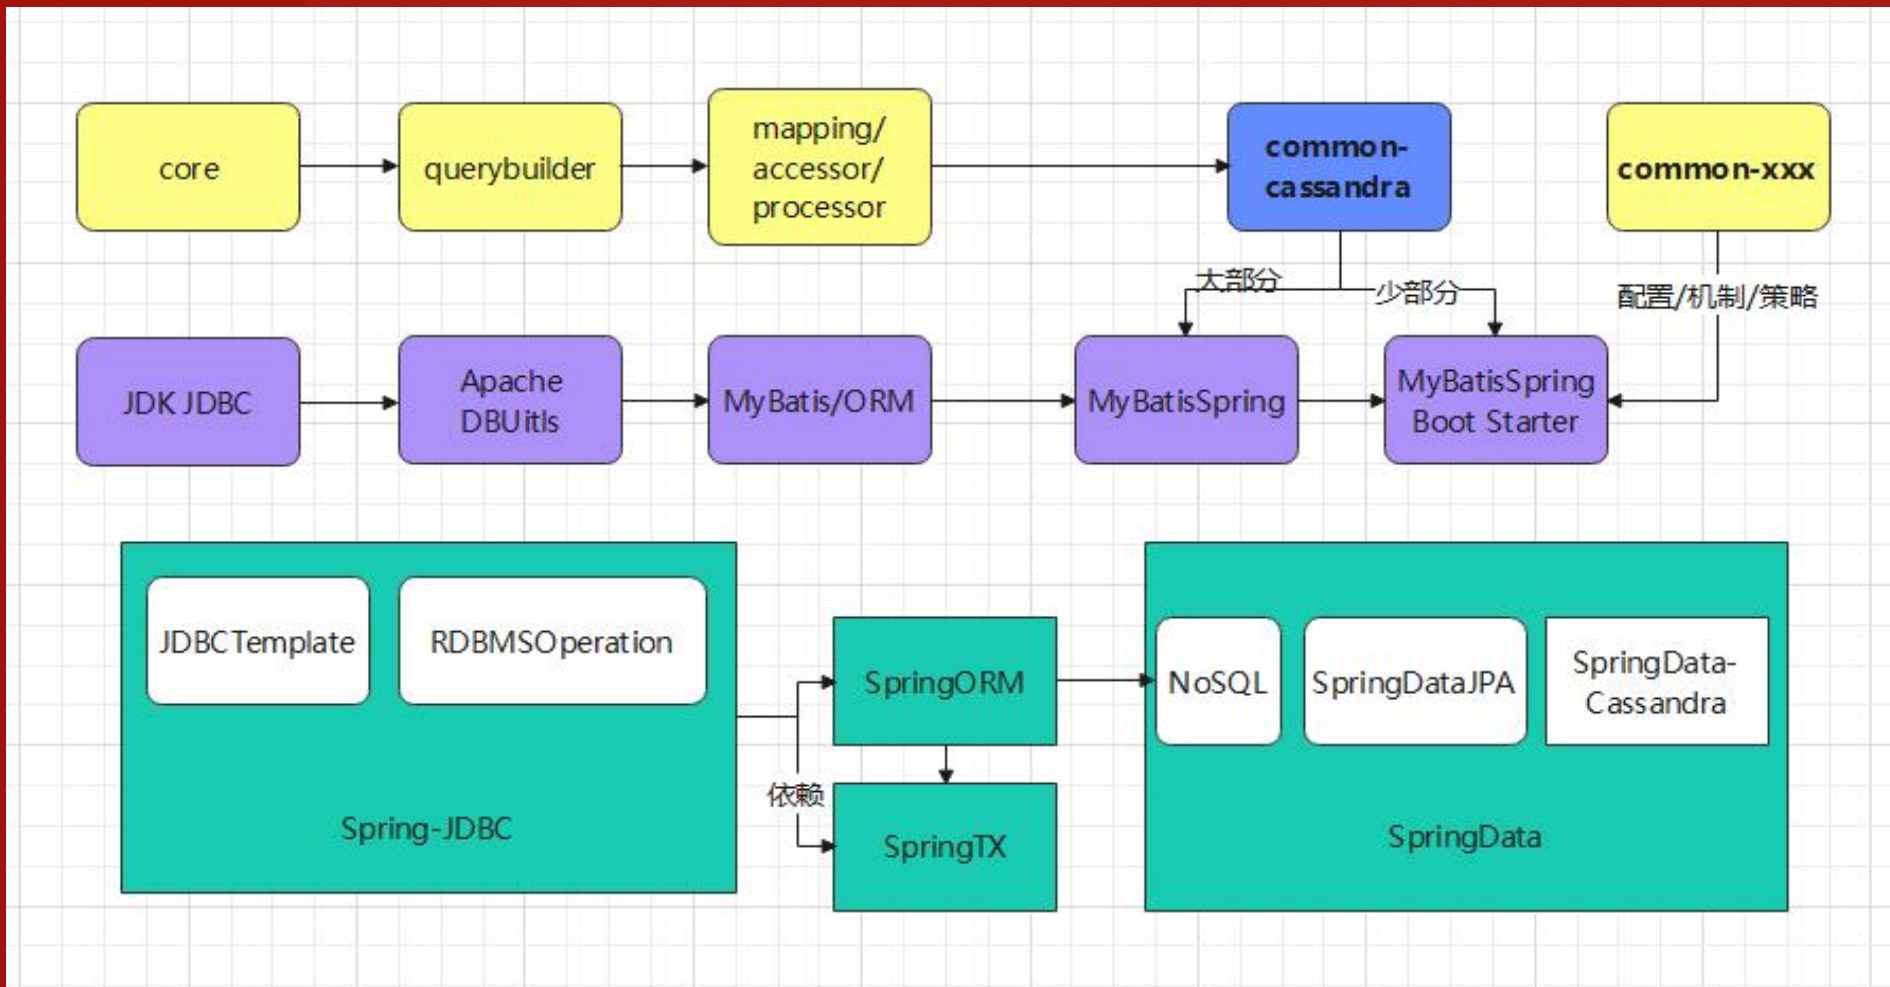
---

- 存储层的封装非常成熟，模式非常固定，机制却因人而异
  - 同步，异步，反应式
  - 技术演进
    - JDBC->DBUtils->Spring JDBC (JdbcTemplate, RDBMSOperation) -> Spring ORM->Spring data.
    - MyBatis->MyBatis-Spring->MyBatis-Spring-starter.
  - 3.11.0 statement/mapper/accessor 4.0.0 processor
- Cluster和Session
  - 所有keyspace一个session
  - 每个keyspace一个session

## 2. 数据建模-客户端

The content of the title

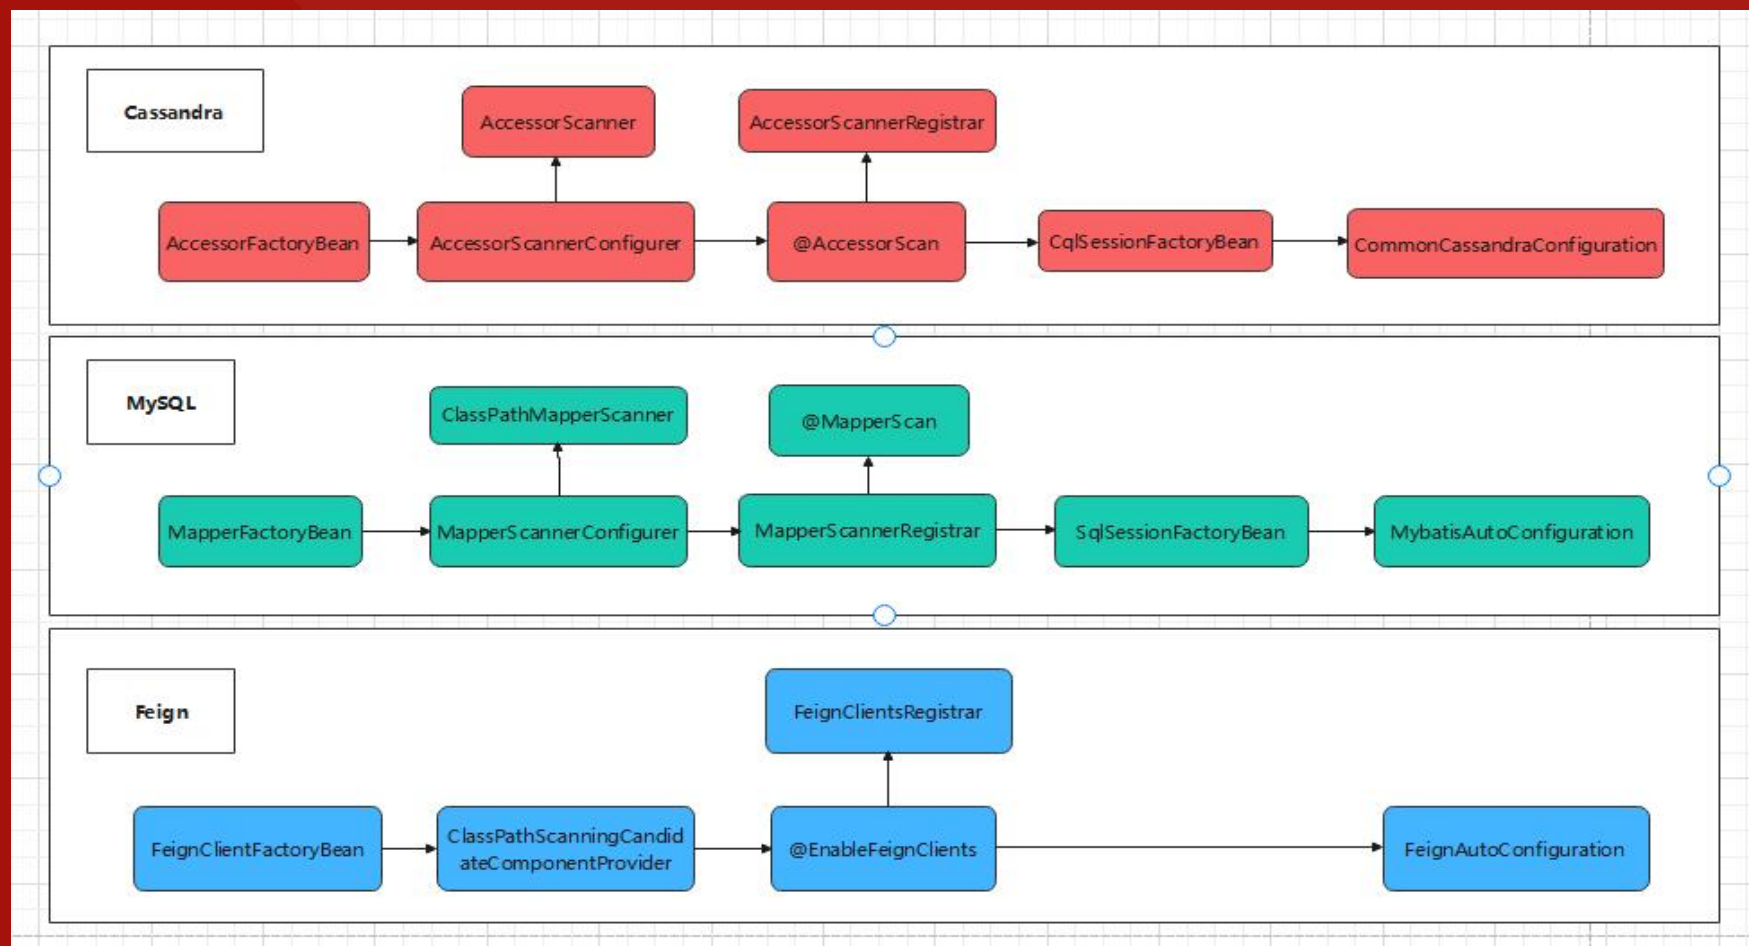
- common-cassandra



## 2. 数据建模-客户端

The content of the title

- common-cassandra





## 2. 数据建模-项目使用

The content of the title

---

- 直播
  - course\_live\_msg: live\_stu\_voice live\_notice live\_all\_chat live\_stu\_praise
- 电商
  - course\_process: user\_lesson\_progress user\_watch\_process user\_point\_history user\_process user\_point\_product
  - course\_naive: user\_profile
- 智学
  - adaplearn\_tiku: answers\_by\_user quiz\_user\_progress progress\_by\_user quiz\_user\_answer exam\_user\_answer
- 题库
  - course\_tiku: user\_answer, user\_live\_answer2, user\_paper\_answer

## 2. 数据建模-CQRS+Event Source

The content of the title

---

- Akka Persistent Cassandra

# 3. 系统架构

The content of the title

---

- 架构
  - 内部实现，外部交互，系统属性
  - 单机特性-对标数据库
    - 数据模型，存储模型，查询模型（读写性能）
  - 分布式特性-对标分布式系统
    - 用户视角：读写路径
    - 系统视角：复制和一致性，容错，分区，共识
  - 其他模型：<https://xiaozhiliaoo.github.io/reading-note/tech-quotes/chapter8.html> 820条

# 3. 系统架构

The content of the title

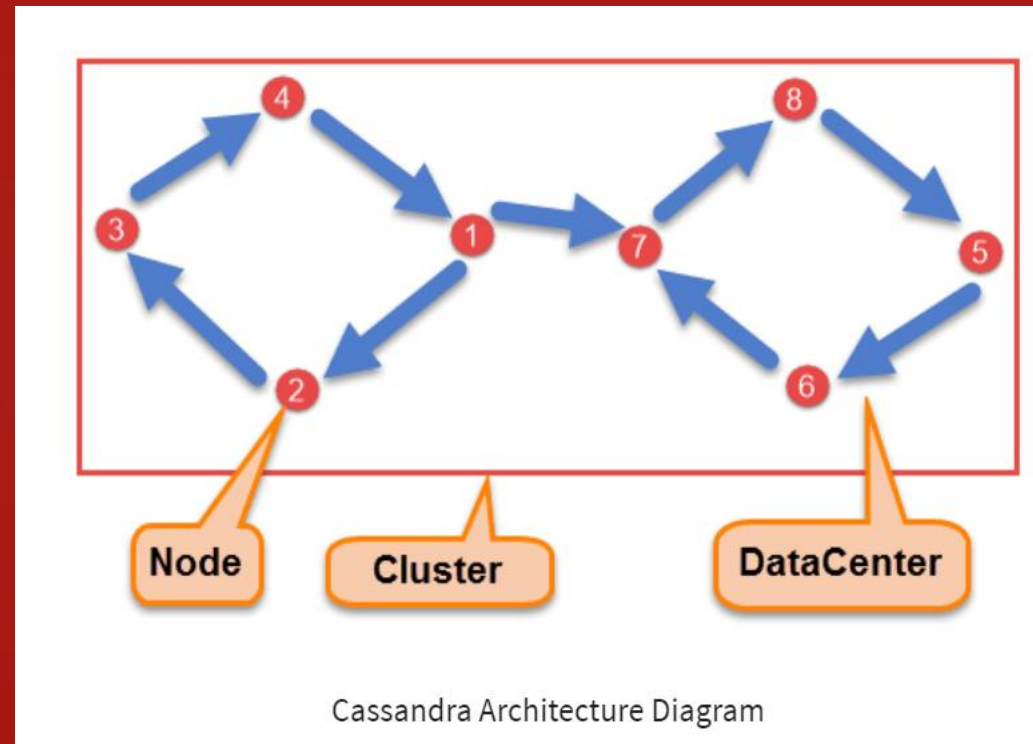
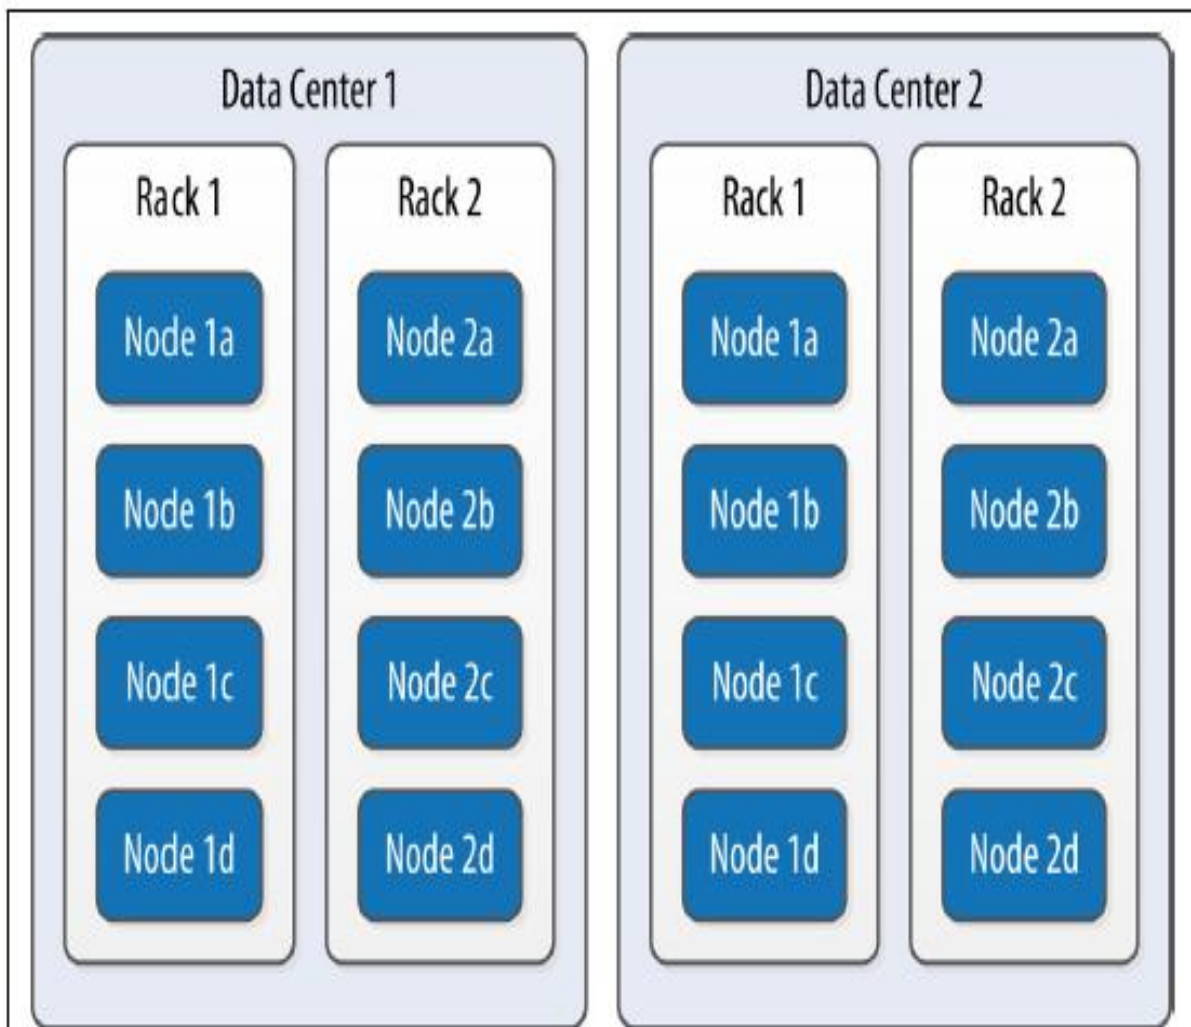
---

- 单机
  - Commit Log, Memmory table, SSTable, 布隆过滤器, 墓碑, 合并, 压缩
- 分布式
  - 数据中心和机架
  - Gossip和故障检测, Snitch感应策略
  - DHT和一致性Hash, 环和令牌, 虚拟节点
  - 分区器
  - 复制策略
  - 一致性级别
  - 查询和协调器节点
  - hint handoff, anti-entory, repair, merkle tree
  - 轻量级事务和Paxos
  - SEDA
- 基本问题
  - 读写路径是什么? (一条SQL执行/打开网页发生了什么)



### 3. 系统架构-数据中心和机架

The content of the title



# 3. 系统架构-通信Gossip, 故障检测, 和Snitch

The content of the title


---

- Gossip
  - Information Dissemination模型的epidemic protocol两大类: anti-entropy和rumor-mongering
  - 设计考虑
    - reconciliation: precise, Scuttlebutt
    - flow control: Spreading Capacity Fairly, Local Adaptation
  - Scuttlebutt: <https://awinterman.github.io/simple-scuttle/>
- 故障检测
  - Phi Accrual FD
- Snitch
  - 确定相对主机远近程度。收集网络拓扑有关信息, 处理路由请求。
  - SimpleSnitch
  - PropertyFileSnitch
  - GossipingPropertyFileSnitch
  - DynamicEndpointSnitch

### 3. 系统架构-Token Ring

The content of the title

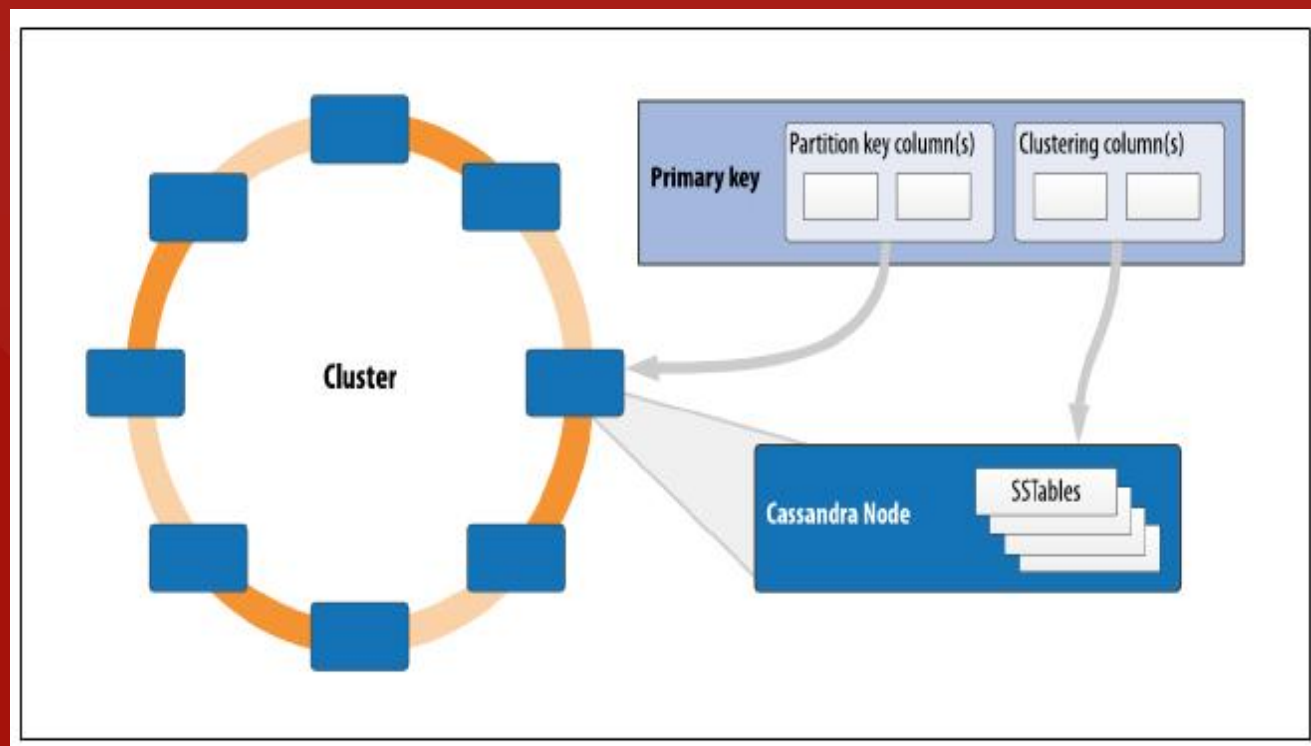
- 一致性Hash和DHT
  - Ring和Range
  - Token
  - 虚拟节点 (num\_tokens)
- Hash算法评估
  - 单调性：旧数据不变性
  - 平衡性：节点各尽所能
    - 平衡不是平均，马太效应
- 通过环保证单调性  
通过虚拟节点保证平衡性



### 3. 系统架构-分区器

The content of the title

- 找到key在token ring的位置,  $\text{token} = \text{partitionFunction}(\text{partition key})$ 
  - Murmur3Partitioner(MP-默认)
    - Mu: multiply R: rotate
    - Redis, Nginx, libmemcached, npm, Hadoop, Cassandra, Solr, Elasticsearch, Guava, Kafka, HBase Lucene
  - OrderPreservingPartitioner(OPP)
  - RandomPartitioner(RP)
  - ByteOrderedPartitioner(BOP)
- Murmur3
  - 即使键是有规律的, 算法仍能给出一个很好的随机分布性, 并且算法的计算速度也非常快



### 3. 系统架构-分区器

The content of the title

#### • 深入Hash算法

##### CRC-64 (1975)

→ Used in networking for error detection.

##### MurmurHash (2008)

→ Designed to a fast, general purpose hash function.

##### Google CityHash (2011)


→ Designed to be faster for short keys (<64 bytes).

##### Facebook XXHash (2012)

→ From the creator of zstd compression.

##### Google FarmHash (2014)

→ Newer version of CityHash with better collision rates.




### 3. 系统架构-复制


The content of the title

- LeaderLess Replication

- Versioned timestamp data(时钟同步), LWW-Element-Set CRDT
- 策略
  - SimpleStrategy
  - NetworkTopologyStrategy
  - LocalStrategy(system keyspace)



SimpleStrategy in Cassandra



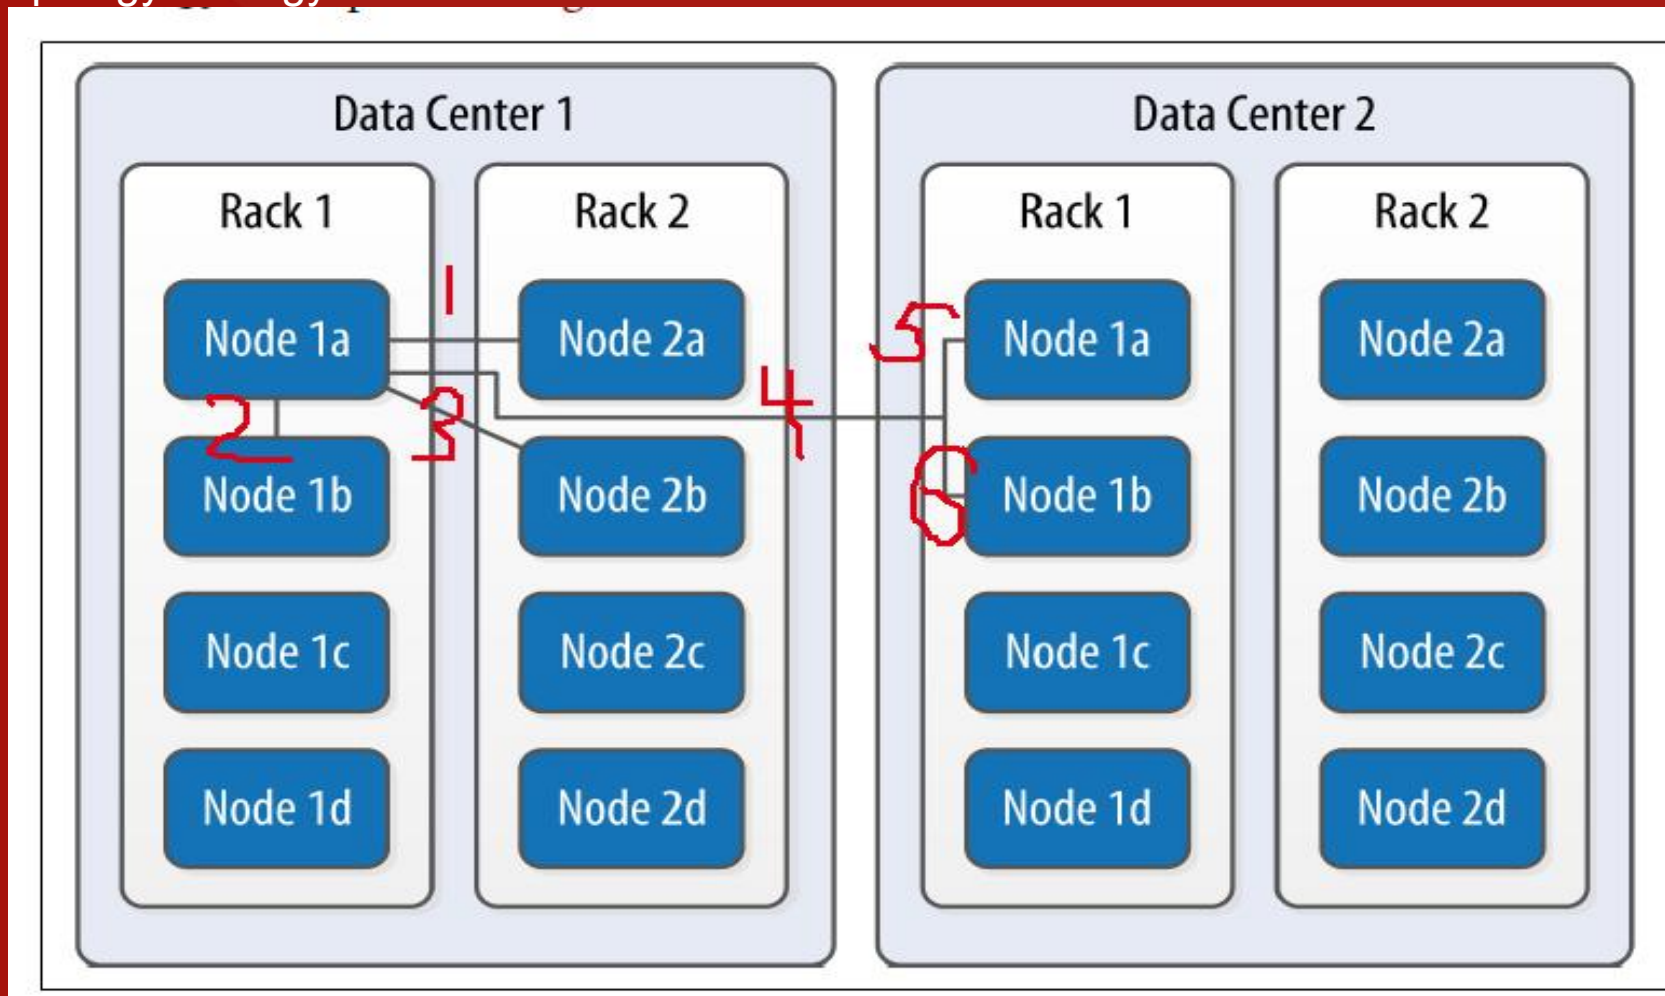
NetworkTopologyStrategy in Cassandra



### 3. 系统架构-复制

The content of the title

- NetworkTopologyStrategy



### 3. 系统架构-一致性

The content of the title

---

- 一致性: Tunable Consistency
  - ANY
  - ONE, TWO, THREE
  - LOCAL\_ONE, QUORUM, LOCAL\_QUORUM, EACH\_QUORUM  
ALL
- Read Path
  - read-repair, Merkle Trees
- Write Path
  - Hinted Handoff(可用性)

### 3. 系统架构-线程模型SEDA

The content of the title

- SEDA: staged event-driven architecture




Figure 6: A SEDA Stage: A stage consists of an incoming event queue, a thread pool, and an application-supplied event handler. The stage's operation is managed by the controller, which adjusts resource allocations and scheduling dynamically.

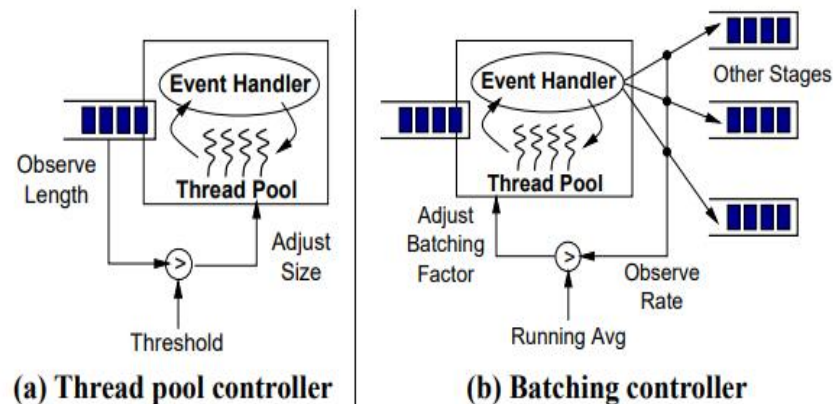


Figure 7: SEDA resource controllers: Each stage has an associated controller that adjusts its resource allocation and behavior to keep the application within its operating regime. The thread pool controller adjusts the number of threads executing within the stage, and the batching controller adjusts the number of events processed by each iteration of the event handler.

### 3. 系统架构-线程模型SEDA

The content of the title

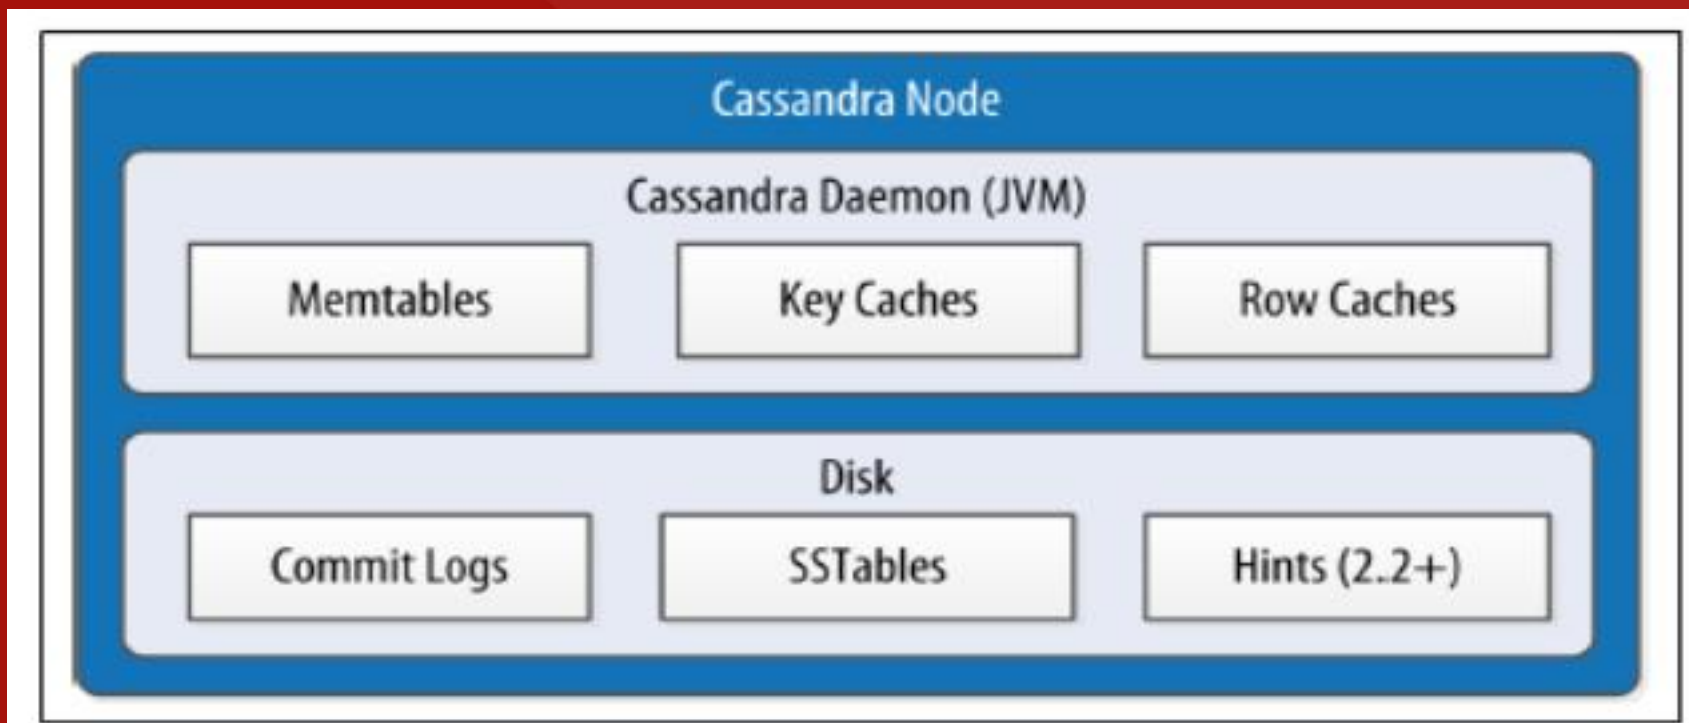
---

- Stage
  - 读，写
  - Gossip
  - 和其他节点交互
  - anti entropy
  - read repair
  - hinted handoff
- org.apache.cassandra.concurrent.Stage

### 3. 系统架构-单机存储引擎

The content of the title

- Memtables, SSTables, and Commit Logs
- Compress
- Compact（读，写，时序）
  - STCS（写多读少），LCS（读多写少），DTCS



### 3. 系统架构-单机存储引擎案例

The content of the title

---


```
CREATE TABLE course_process.user_process (  
    userid text,  
    courseid text,  
    lessonid text,  
    all bigint,  
    duration bigint,  
    finish bigint,  
    mark bigint,  
    time bigint,  
    PRIMARY KEY (userid, courseid, lessonid)  
) WITH CLUSTERING ORDER BY (courseid ASC, lessonid ASC)  
    AND bloom_filter_fp_chance = 0.01  
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}  
    AND comment = ''  
    AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}  
    AND compression = {'chunk_length_in_kb': '64', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}  
    AND crc_check_chance = 1.0  
    AND dclocal_read_repair_chance = 0.1  
    AND default_time_to_live = 0  
    AND gc_grace_seconds = 864000  
    AND max_index_interval = 2048  
    AND memtable_flush_period_in_ms = 0  
    AND min_index_interval = 128  
    AND read_repair_chance = 0.0  
    AND speculative_retry = '99PERCENTILE';  
CREATE INDEX lesson_index ON course_process.user_process (lessonid);
```



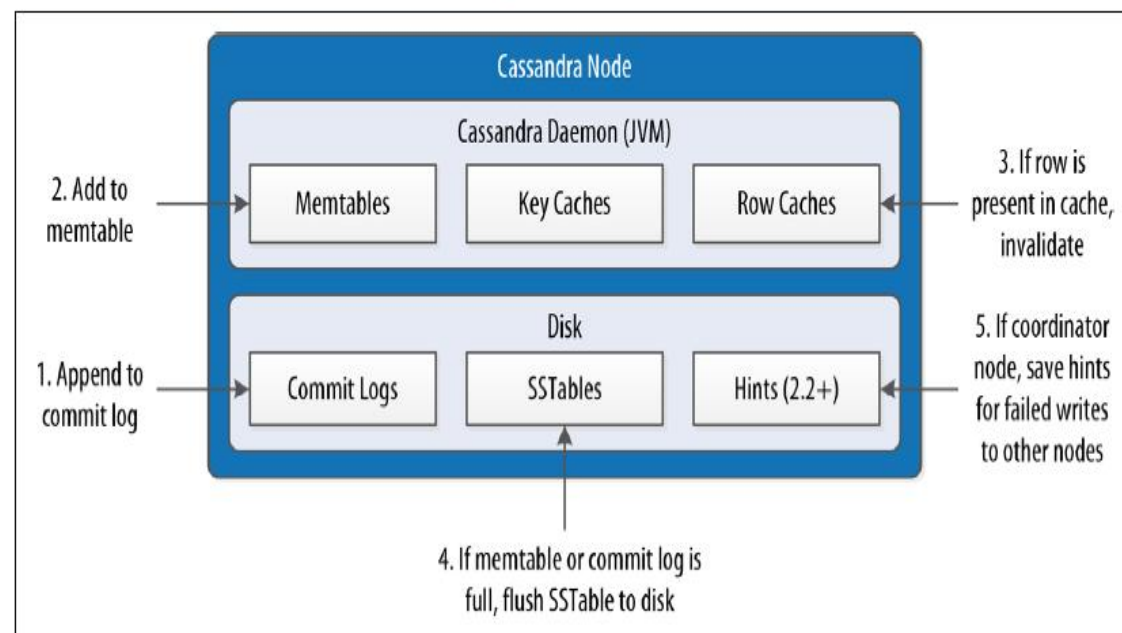
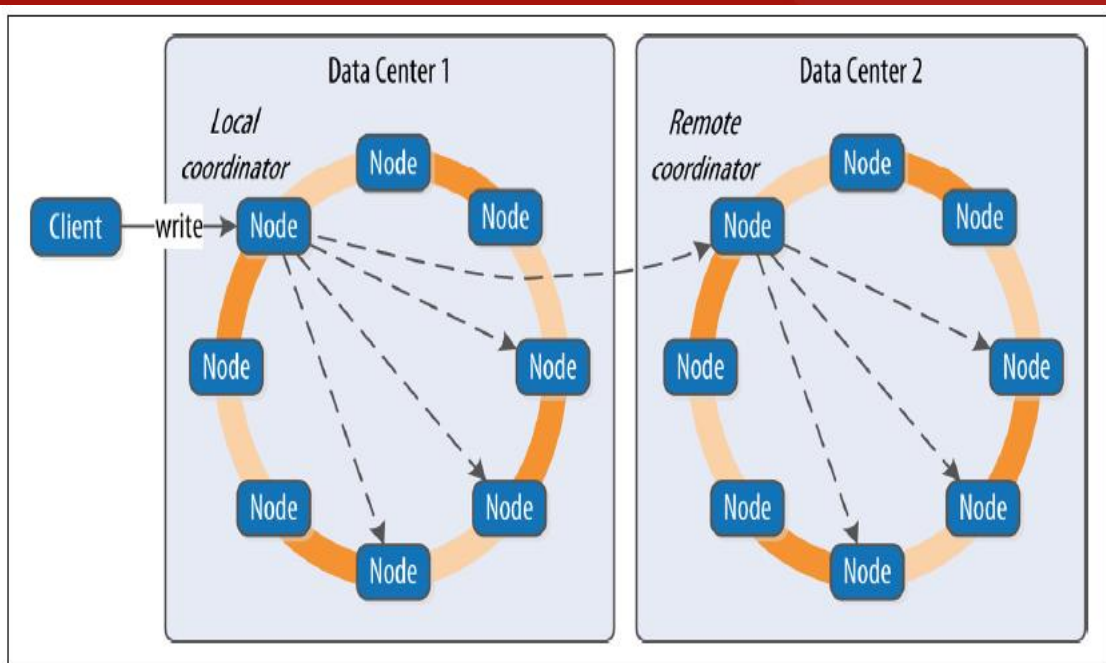
### 3. 系统架构-写路径

The content of the title

- 正常写
- 轻量级事务(条件插入, 更新)
- 批处理
- 删除



Write Operation in Cassandra



### 3. 系统架构-读路径

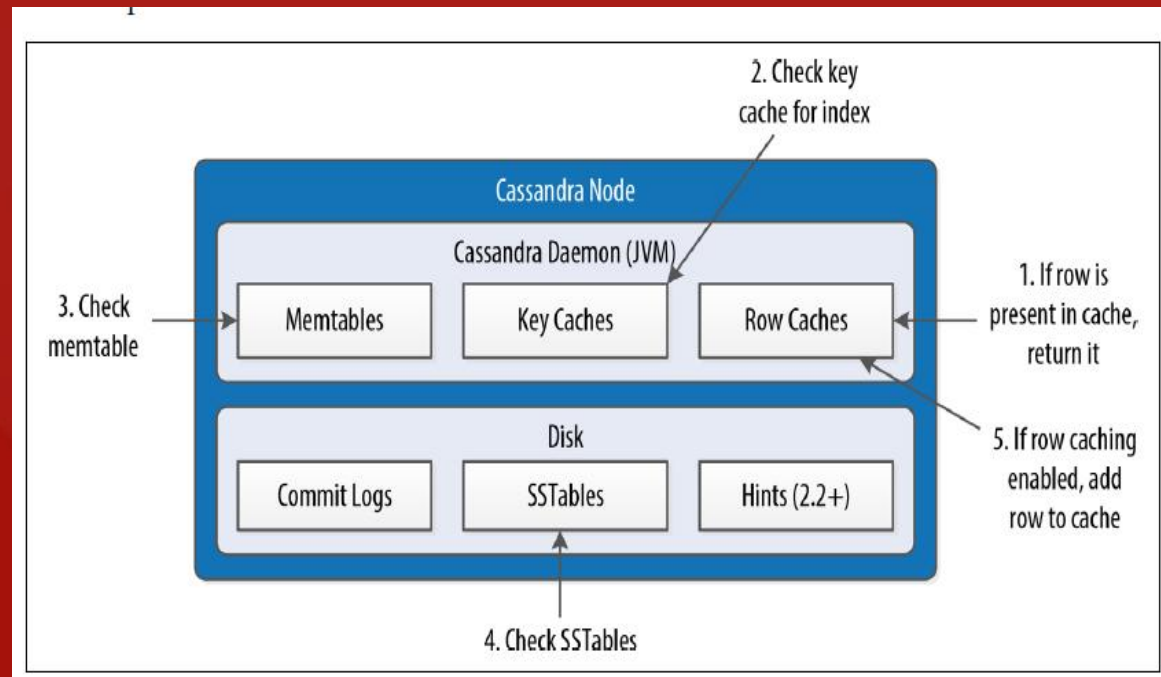
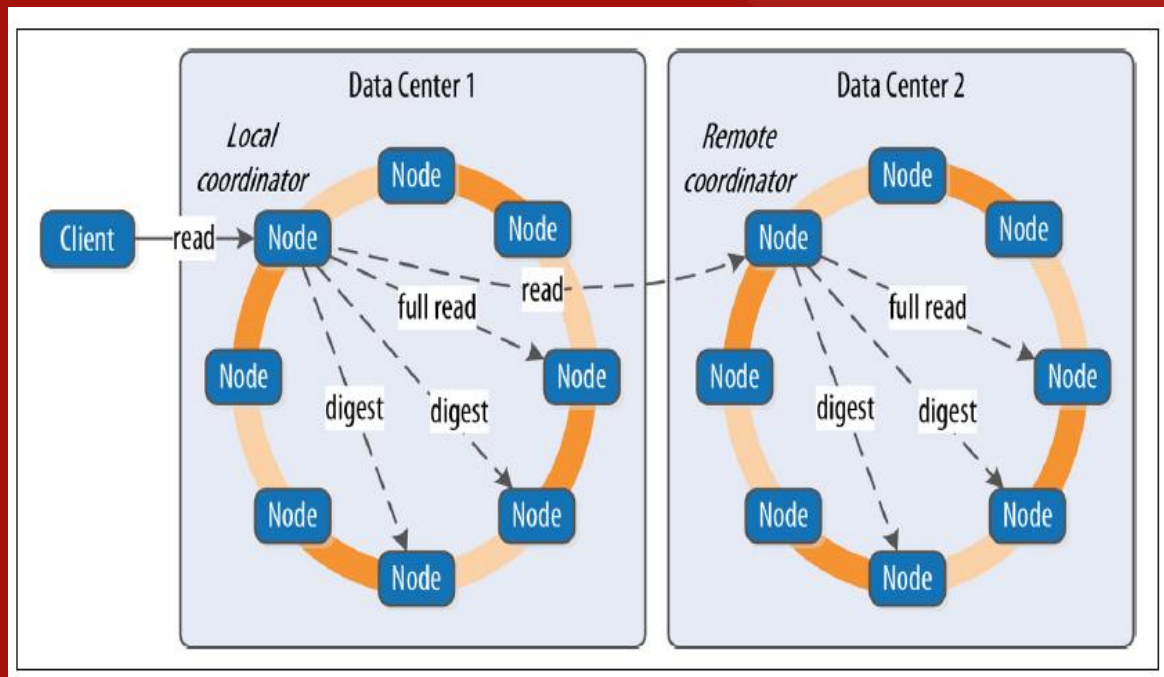
The content of the title

- 正常读/allow filtering查询
- 范围查询
- 排序, 过滤, 分页, UDF, UDA
- 二级索引, 物化视图

```
... PRIMARY KEY ((hotel_id, start_date), room_number)
... ) WITH comment = 'Q7. Find reservations by hotel and date';
cqlsh:adaplearn_tiku_test>
cqlsh:adaplearn_tiku_test> CREATE MATERIALIZED VIEW reservations_by_confirmation AS
... SELECT * FROM reservations_by_hotel_date
... WHERE confirm_number IS NOT NULL and hotel_id IS NOT NULL and
... start_date IS NOT NULL and room_number IS NOT NULL
... PRIMARY KEY (confirm_number, hotel_id, start_date, room_number);
```

Warnings :

Materialized views are experimental and are not recommended for production use.



## 4. NoSQL存储选型

The content of the title

- 模型：属性占比法，决策树，案例分析法，成本收益分析法
- 如何比较存储？
  - 单机 ---> 最终会对标单机数据库
  - 分布式 ---> 最终会对标CAP/PACELC
- 同类优先性能，不同优先功能
  - 性能提高不易，功能实现变易。

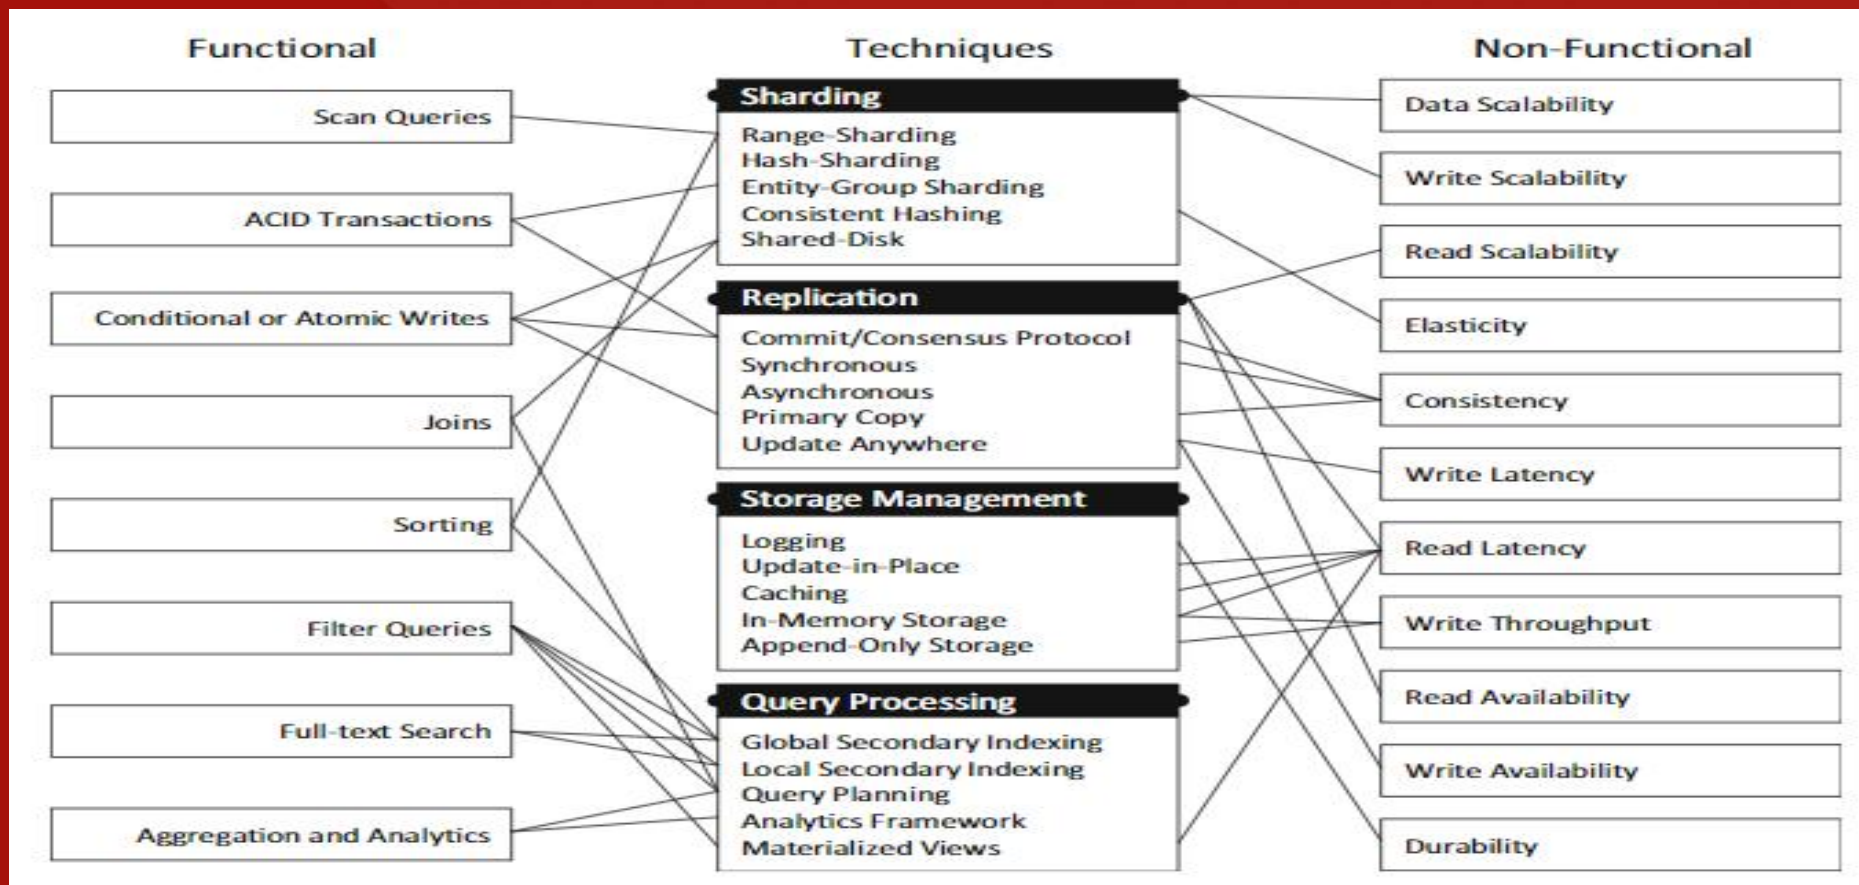
| 行业常用MQ对比 |         |              |           |             |                 |
|----------|---------|--------------|-----------|-------------|-----------------|
| 分类       | 对比项     | RabbitMQ     | RocketMQ  | Kafka       | Pulsar          |
| 基础能力     | 架构      | 单体           | 单体        | 单体          | 计算存储分离          |
|          | 消费模式    | 推 + 拉        | 推 + 拉     | 拉模式         | 推 + 拉           |
|          | 重试队列    | 不支持          | 支持        | 不支持         | 支持              |
|          | 死信队列    | 支持           | 支持        | 不支持         | 支持              |
|          | 海量Topic | 不支持          | 支持        | 不支持         | 支持              |
|          | 延迟消息    | 支持           | 支持        | 不支持         | 支持              |
|          | 堆积能力    | 一般           | 海量        | 海量          | 海量              |
|          | 消息回溯    | 不支持          | 支持        | 支持          | 支持              |
|          | 多协议支持   | AMQP、MQTT等协议 | 私有协议      | 私有协议        | 私有协议、AMQP、MQTT等 |
|          | 安全机制    | 支持           | 支持        | 支持          | 支持              |
| 服务能力     | 事务性消息   | 支持           | 支持        | 支持          | 支持              |
|          | 吞吐量     | 一般           | 高         | 非常高         | 非常高             |
|          | 低延迟     | 非常好          | 好         | 一般(分区数多越明显) | 好               |
|          | 可靠性     | 主备模式         | 多副本同/异步刷盘 | 多副本异步刷盘     | 多副本同/异步刷盘       |
|          | 一致性     | 主从模式         | 主从模式      | ISR算法       | Quorum算法        |
| 运营能力     | 可用性     | 一般           | 较高        | 较高          | 高               |
|          | 多租户     | 支持           | 不支持       | 不支持         | 支持              |
|          | 动态扩容    | 横向扩容         | 需手动同步配置   | 需平衡数据       | 友好(即时扩容)        |
|          | 故障恢复    | 不友好          | 不友好       | 较友好         | 友好              |
|          | 数据清理    | Topic级别      | 集群级别      | Topic级别     | Topic级别         |
| 使用场景     | 安全机制    | 身份认证+权限      | 不支持       | 身份认证+权限     | 身份认证+权限         |
|          | 典型场景    | 传统业务场景       | 高性能电商场景   | 高吞吐大数据场景    | 金融场景和大数据场景      |



## 4. NoSQL存储选型

The content of the title


- 属性分析法-NoSQL Toolbox：你觉得哪些因素最重要？



## 4. NoSQL存储选型

The content of the title

- 决策树模型：你觉得哪些选择很重要？



## 4. NoSQL存储选型

The content of the title

---

- 技术工具：<https://dbdb.io/>
  - 查找竞品
    - Hazelcast竞品
  - 研究学习
    - 学习Java，MVCC，嵌入式，开源
  - 根据特性查找
    - MapDB



## 5. 参考资料

The content of the title

---

- 项目
  - <https://github.com/xiaozhiliaoo/cassandra-practice>
- 书
  - 《Cassandra权威指南-第二和三版》
  - 《数据库系统内幕》 第7,9,11,12章
  - 《数据密集型应用系统设计》 第2,3,5,6,9章
  - 《云计算与分布式系统》 第8章
- 论文
  - Dynamo/Bigtable
  - 原始论文: Cassandra - A Decentralized Structured Storage System
  - 线程模型: SEDA: An Architecture for Well-Conditioned, Scalable Internet Services
  - 通信协议: Efficient Reconciliation and Flow Control for Anti-Entropy Protocols
  - 数据建模: A Big Data Modeling Methodology for Apache Cassandra
  - NoSQL选型: NoSQL database systems: a survey and decision guidance

# 6. 回顾

The content of the title

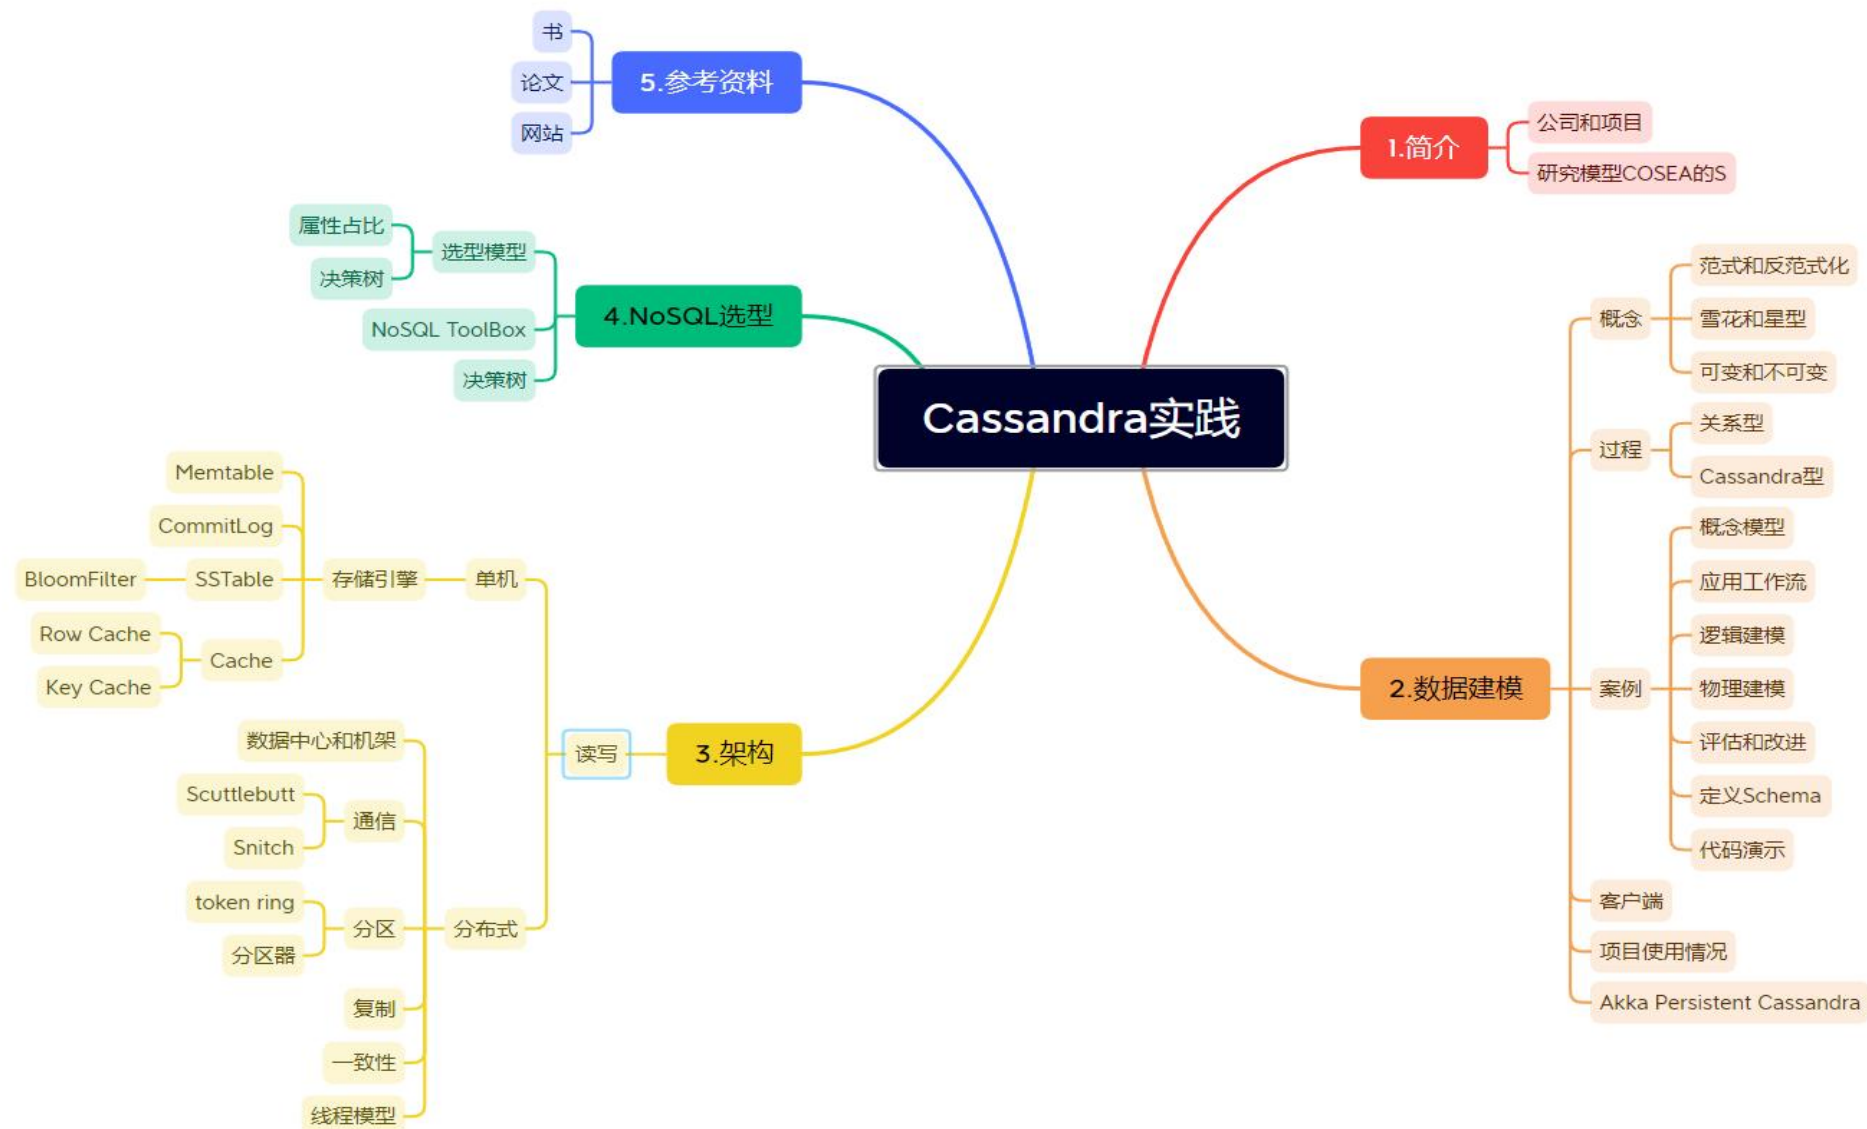
---

## • 分布式系统泛型

- 体系结构：非集中式（去中心化）
- 进程：SEDA
- 通信：多播-Gossip(逆熵, Scuttlebutt), Snitch, Information 【Vector clock(generation(monotonic timestamp), version(logical clock))】
- 命名：无层次命名-DHT
- 协调：时钟同步, Last-Write-Wins-Element-Set, Paxos-LWT
- 一致性和复制：可调一致性, 逆熵, 修复, 一致性Hash复制
- 容错性：FD故障检测器
- 安全：认证, 授权, 加密, SSL, TLS和证书

# 6. 回顾

The content of the title



The background features a complex, abstract architectural line drawing in a lighter shade of red. The drawing consists of various geometric shapes, including rectangles, triangles, and trapezoids, which are interconnected to form a sense of depth and perspective. The lines are thin and precise, creating a modern, architectural feel. The overall composition is centered around the text, with the drawing elements framing it from the sides and bottom.

# Q&A

網易 NETEASE



# THANK YOU

網易 NETEASE